

Quantitative Methods 2, ZHAW

Jürgen Degenfellner

2025-02-17

Contents

| | |
|--|-----------|
| 1 Preamble | 5 |
| 1.1 Books we will heavily borrow from are: | 5 |
| 1.2 If you need a good reason to buy great books... | 6 |
| 2 Introduction | 7 |
| 2.1 What is statistical modeling and what do we need this for? | 7 |
| 2.2 A (simple) model for adult body heights in the Bayesian framework | 11 |
| 2.3 Classical approach for the simplest model | 21 |
| 2.4 Exercises | 27 |
| 2.5 Addendum | 28 |
| 3 Simple Linear Regression | 33 |
| 3.1 Simple Linear Regression in the Bayesian Framework | 33 |
| 3.2 Simple Linear Regression in the Frequentist Framework | 47 |
| 3.3 Exercises | 71 |
| 4 Multiple Linear Regression | 75 |
| 4.1 Linear Regression with 2 predictors in the Bayesian Framework . | 75 |
| 4.2 Linear regression with 2 predictors in the Frequentist Framework | 84 |
| 4.3 Exercises | 105 |
| 4.4 TODOS | 108 |

Chapter 1

Preamble

This script is a continuation of the first one for Quantitative Methods 1 at ZHAW.

In the first part, we learned about the basics of probability theory, descriptive statistics, Bayesian statistics, and hypothesis testing.

In this script, we will dive into the basics of statistical modeling - a world of aesthetic wonder and surprises.

This script is a first draft as you are the first group to be working with it.

Please feel free to send me suggestions for improvements or corrections.

This **should be a collaborative effort** and will (hopefully) never be finished as our insight grows over time.

The script can also be seen as a pointer to great sources which are suited to deepen your understanding of the topics. Knowledge is decentralized, and there are many great resources out there.

For the working setup with R, please see this and the following sections in the first script.

The complete code for this script can be found [here](#).

1.1 Books we will heavily borrow from are:

- (older online version is Free; current version in ZHAW Library) Statistical Rethinking, YouTube-Playlist: Statistical Rethinking 2023
- (Free) Understanding Regression Analysis: A Conditional Distribution Approach

- (Free online-access via ZHAW Library) Data Analysis Using Regression and Multilevel/Hierarchical Models
- (Free) Doing Bayesian Data Analysis



1.2 If you need a good reason to buy great books...

Think of the total costs of your education. You want to extract maximum benefit from it. In the US, an education costs a lot. In beautiful Switzerland, the tuition fees (if applicable) are nowhere near these figures. Costs you could consider are opportunity costs of not working. A comparison with both, a foreign education or opportunity costs, justifies the investment in good books. Or: The costs of all the good books of your education combined are probably less than an iPhone Pro.

Chapter 2

Introduction

2.1 What is statistical modeling and what do we need this for?

Typically, one simplifies the complex reality (and loses information) in order to make it better **understandable** (explainable), mathematically treatable and to make **predictions**.

Underlying our models, there are theories which should be falsifiable and testable. For instance, I would be really surprised if I pull up my multimeter and measure the voltage (V) and electric current (I) at a resistance (R) in a circuit and find that Ohm's law $V = IR$ is not true. This **law** can be tested over and over again and if one would find a single valid counterexample, the law would be falsified. It is also true that the law is probably not 100% accurate, but an extremely precise approximation of reality. Real-world measurements carry measurement errors and when plotting the data, one would see that the data points might not lie exactly on a straight line. This is not a problem.

A statistical model is a mathematical framework that represents the relationships between variables, helping us understand, infer, and predict patterns in data. It acts as a bridge between observed data and the real-world processes that generated them. In health research, where variability and uncertainty are inherent, statistical models are valuable tools for making sense of complex phenomena. You can watch this as short intro.

In QM1 we have already made testable predictions with respect to the probability of an event. In our 1000-researcher experiment we stated for instance, that the probability of observing 66 or more findings would be very unlikely. If such an event would occur (while not repeating the experiment many times), we would reconsider our model. Inexactly, we could have stated something like:

“We will not see more than 100 findings by chance.” With respect to our multiple choice test at the end of QM1 we could predict: “We will not see a single person answering all questions correctly by chance in our lifetime (given the frequency of tests).” Note, that in this context, the word **predict** is used with respect to a future event (chance finding or chance passing of the test). As we will see, there does not necessarily have to be a temporal connection in order to *predict* something.

Depending on the task at hand, we would use different models. In any case, logical reasoning and critical thinking comes first, then comes the model. **It makes no sense to estimate statistical models just for the sake of it.**

All models are wrong, but some are useful. Or to quote George Box:

“Since all models are wrong the scientist cannot obtain a ‘correct’ one by excessive elaboration. On the contrary following William of Occam he should seek an economical description of natural phenomena. Just as the ability to devise simple but evocative models is the signature of the great scientist so overelaboration and overparameterization is often the mark of mediocrity.”

In my opinion, statistical modeling is an art form: difficult and beautiful.

One goal of this course is to improve interpretation and limitations of statistical models. They are not magical tools turning data into truth. Firstly, the rule garbage in, garbage out (GABA) applies. Secondly, statistical models are based on data and their variability and have inherent limitations one cannot overcome even with the most sophisticated models. This is expressed for instance in the so-called bias-variance trade-off. You can’t have it all.

2.1.1 Explanatory vs. Predictive Models

I can recommend reading this article by Shmueli et al. (2010) on this topic.

Statistical models serve different purposes depending on the research question. Two primary goals are **explanation** and **prediction**, and each requires a different approach:

Explanatory models (the harder of the two) focus on understanding causal relationships. These models aim to uncover mechanisms and answer “**why**” questions. For example:

- Does smoking increase the risk of lung cancer? **Yes.** (If you want to see what a large effect-size looks like, check out this study.)
- How large is the **effect** (causal) of smoking on lung cancer? **Large.**
- Does pain education and graded sensorimotor relearning improve disability (a question we ask in our Resolve Swiss project)?

2.1. WHAT IS STATISTICAL MODELING AND WHAT DO WE NEED THIS FOR?9

Explanatory models are **theory-driven**, designed to test hypotheses. Here, one wants to understand the underlying mechanisms and the relationships between variables and hence often uses (parsimonious) models that are more interpretable, like linear regression.

Predictive models prioritize forecasting/predicting (future) outcomes based on patterns in the data. These models aim to answer “**what will happen?**” For instance:

- Gait analysis using Machine Learning (ML)?
- Skin cancer detection using neural networks?
- If I know the age, sex and weight of a person, can I predict his/her height?
Can I predict the height better with the given covariates compared to just guessing by using the mean height of my sample for the next patient?

Predictive models are **data-driven**, often using complex algorithms to achieve high accuracy. Their success is measured using metrics like Root Means Square Error (RMSE), Area Under the Curve (AUC), or **prediction error on new, unseen data**. Any amount of model complexity is allowed. One could for instance estimate a neural network (“just” another statistical model) with many hidden layers and neurons in order to improve prediction quality. Interpretability of the model weights is not a priority here.

While explanatory and predictive goals often complement each other, their differences highlight the importance of clearly defining the purpose of your analysis. In applied health research, explanatory models help identify causal mechanisms, while predictive models can guide real-world decisions by providing actionable forecasts. Together, they enhance both our understanding of phenomena and our ability to make informed decisions in complex environments.

2.1.2 Individual vs. Population Prediction

Another important distinction is between **individual vs. population** prediction. In the smoking example above, we can be very sure about the mean effects that smoking has on lung cancer. On an individual level, it is harder to predict the outcome. Nevertheless, individual predictions will be (notably) better than random guessing. We will discuss this in greater detail.

Statistical models are often much worse than one would naively expect, but they very often better than experts. If you are interested and want to boost your confidence in the predictive ability of statistical models, I recommend reading chapter 21 (“Intuitions vs. Formulas”) of Daniel Kahneman’s book “Thinking, Fast and Slow” (available in the ZHAW library).

2.1.3 Practical Use of Statistical Models

In my opinion, we should never be afraid to test our statistical models (as honestly as possible) against reality. We could for instance ask ourselves:

- “How much better does this model estimate and outcome than the arithmetic mean? (i.e., the linear model with just an intercept)”
- “How much better does this model classify than random guessing?”
- Is it worth the effort to collect data and estimate this model by using hundreds of hours of our time?

In some cases, these questions can be answered straightforwardly.

- In advertising (Google, Facebook, ...), a couple of percentage points in prediction quality might make a difference of millions of dollars in revenue offsetting the statistitians salary by a large margin.
- Improved forecasts of a few percentage points in the stock market or just being slightly better than the average, will make you faboulously rich.
- Improved cancer forecasting might save lives, money and pain and is of course not only measured in financial gains.

2.1.4 Start at the beginning

What do we actually want to do in general? Very broadly speaking we want to: **describe** the association of variables to each other that carry variability. Hence, the relationship is not deterministic like

$$y = 2x + 3$$

but rather we need to “loosen up” the relationship to account for variability (in x and y). So, y and x are not fixed but afflicted with uncertainty. Depending on your philosophical view, you might say you want to find the “true” but unknown relationship (here, 2 and 3 are the true coefficients) between variables. This is what we do in simulation studies all the time: We know the true relationship, simulate data by adding variability and then try to estimate the true relationship we assumed in the first place. This is an **advantage** the pioneers of statistics did not have. We can simulate millions of lines of data at the click of a button. For some practical applications, we can get a really nice and complete answer to our question (for instance sample size for proportions).

So we are looking for a function f such that

$$\mathbf{Y} = f(\mathbf{X})$$

where

- \mathbf{Y} is the “outcome”, “dependent variable” or “response”.
- \mathbf{X} are the “predictors”. \mathbf{X} can be a single Variable x or many variables x_1, x_2, \dots, x_p .

It is important to be aware of the notation here: “Predict” does **not necessarily** mean that we can predict the value in the future. It merely means we estimate the value (or mean) of Y given X .

- This can be done at the same time points, known as **cross-sectional** analysis (“What is the maximum jumping height of a person given their age at a certain point in time, whereas both variables are measured at the same time?”);
- or at different time points, known as **longitudinal analysis** (“What is the maximum jumping height of a person 10 years later (t_2) given their baseline health status at time t_1 ?”).

The **simplest statistical model** would be the **mean model** where Y is “predicted” by a constant: $Y = c$ which (at least in the classical linear regression) turns out to be $c = \bar{x}$. This simple model is often surprisingly good, or, to put it in other words, models with more complexity are often not that much better.

2.2 A (simple) model for adult body heights in the Bayesian framework

As repetition, read the parts about Bayes statistics from QM1 again to refresh your memory about the Bayesian framework.

It's recommendable to read the beginning of the book Statistical rethinking (hint: the online-version of the book differs a bit from the paper-version) up until page 39 as well. We are not completely new to the topic of Bayes thanks to QM1. In the Bayesian setting we use (well argued for) prior knowledge about a parameter or effect and update this knowledge with new data.

We want to **start building our first model** right away.

Let's begin with the example in Statistical rethinking using data from the !Kung San people.

```

library(rethinking)
data("Howell1")
d <- Howell1
str(d)

## 'data.frame': 544 obs. of 4 variables:
## $ height: num 152 140 137 157 145 ...
## $ weight: num 47.8 36.5 31.9 53 41.3 ...
## $ age   : num 63 63 65 41 51 35 32 27 19 54 ...
## $ male  : int 1 0 0 1 0 1 0 1 0 1 ...

d2 <- d[d$age >= 18, ] # only adults

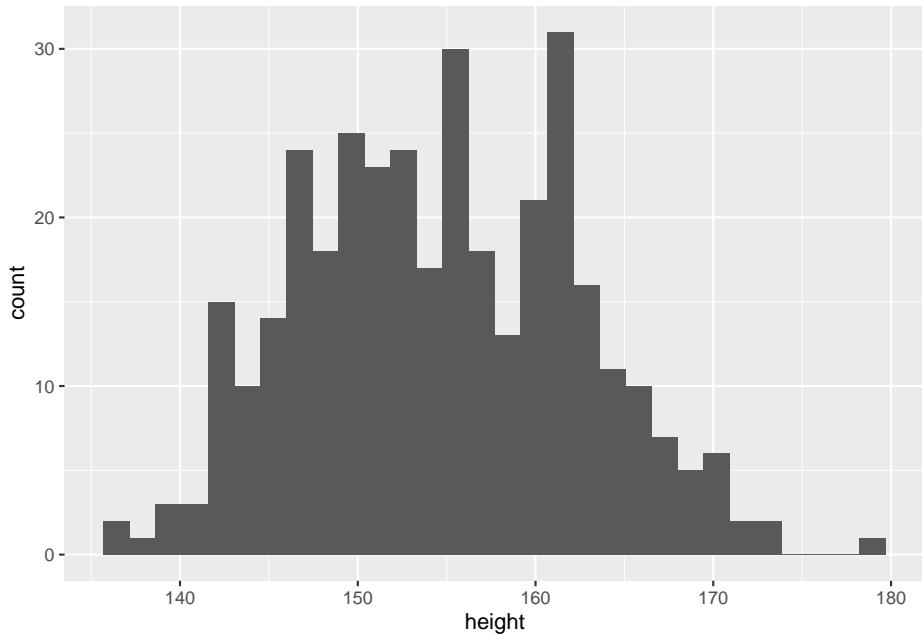
```

We want to model the adult height of the !Kun San people using prior knowledge (about the Swiss population) and data.

```

library(tidyverse)
d2 %>% ggplot(aes(x = height)) + geom_histogram()

```



Since we already have domain knowledge in this area, we can say that heights are usually normally distributed, or at least a mixture of normal distributions (female/male). We assume the following model:

$$h_i \sim \text{Normal}(\mu, \sigma)$$

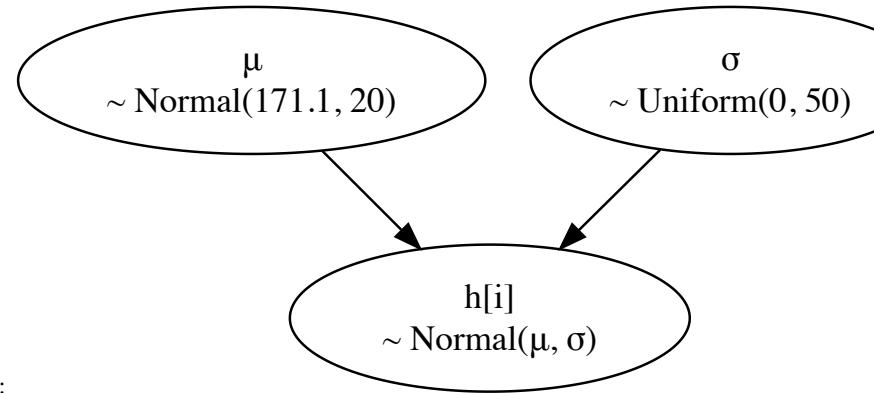
As in QM1, we want to start with a Bayesian model and hence, we need some priors.

Since we are in Switzerland and just for fun, we use the mean of Swiss body heights as expected value for the **prior for the mean**. According to the link (Bundesamt für Statistik), the mean height of $n = 21,873$ people in the Swiss sample is 171.1 cm. We choose the same σ for the prior of the normal as in the book not to deviate too much from the example at hand.

Next comes our **model definition in the Bayesian framework**, which I often find more intuitive than the Frequentist approach:

$$\begin{aligned} h_i &\sim \text{Normal}(\mu, \sigma) \\ \mu &\sim \text{Normal}(171.1, 20) \\ \sigma &\sim \text{Uniform}(0, 50) \end{aligned}$$

Description of the model definition: The heights are normally distributed with unknown mean and standard deviation. As our current knowledge about the mean height, we use a prior distribution for the mean (we do not know but want to estimate) by assuming the mean of a population we know and a standard deviation of 20 cm which allows a rather large range of possible values for μ (the unobserved population mean of the !Kung San people). σ (the unobserved standard deviation of the population of !Kun San people) is also unknown and a priori we restrict ourselves to values between 0 and 50 cm, whereas we assign equal plausibility to all values in this range (which can and should be critically discussed).



Vizualisation of the model structure:

Mind that there is a **conceptual difference** between the normal distribution of the heights and the normal prior distribution of the mean. The latter expresses our prior knowledge/insecurity about the unobserved mean of the normal distribution of the heights. The normal distribution of the heights says we expect the heights to be normally distributed but we do not know the parameters (μ)

and σ) yet. We will estimate these parameters using prior knowledge and the data.

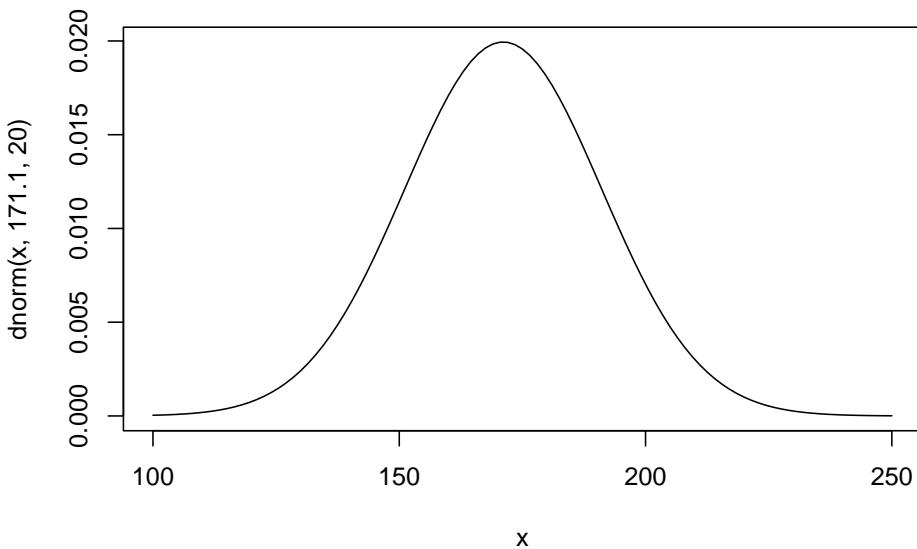
Of course we would not need the prior here due to the large sample size, but let's do it anyways for demonstration purposes. We are not completely uninformed about body heights and express our knowledge with the prior for μ . The 20 in the prior for the mean expresses our range of possible true mean values and acknowledge that there are a variety of different subpopulations with different means.

Using the Swiss data in the link one could estimate that the standard deviation of the heights from 21,873 Swiss people is around is 25.67 cm (Exercise 1).

Remember, in the Bayesian world, there is no **fixed but unknown** parameter, but instead we define a distribution over the unobserved parameter.

We **visualize the prior for μ** :

```
curve(dnorm(x, 171.1, 20), from = 100, to = 250)
```

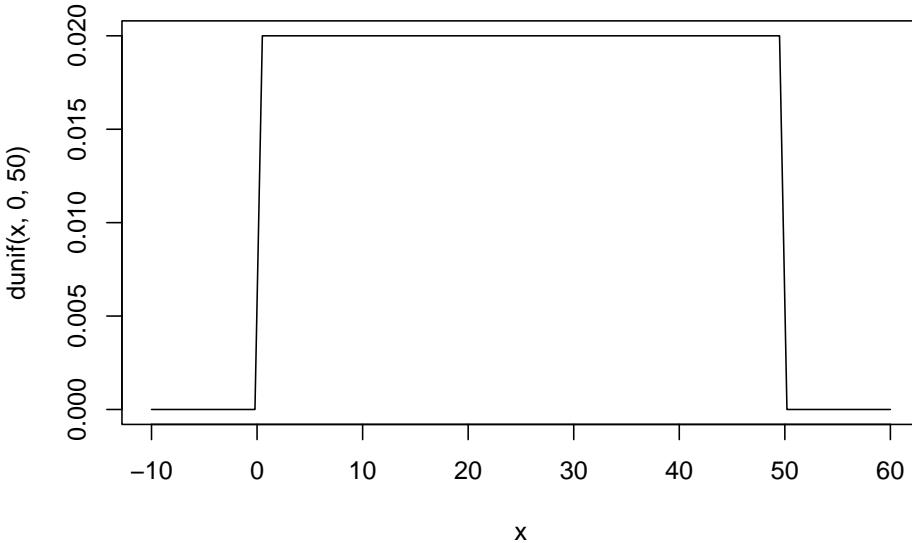


A wide range of population means is possible. Once could discuss this distribution and maybe further restrict it.

The **prior for σ** is uniform between 0 and 50 cm. This is a very wide prior and just constrains the values to be positive and below 50 cm. This could be stronger of course.

Visualization of the prior for σ :

```
curve(dunif(x, 0, 50), from = -10, to = 60)
```



Note, we didn't specify a prior probability distribution of heights directly, but once we've chosen priors for μ and σ , these imply a prior distribution of individual heights.

Without even having seen the **new data**, we can check what our prior (model) for heights would predict. This is important. If the prior already predicts impossible values, we should reconsider our priors and/or model.

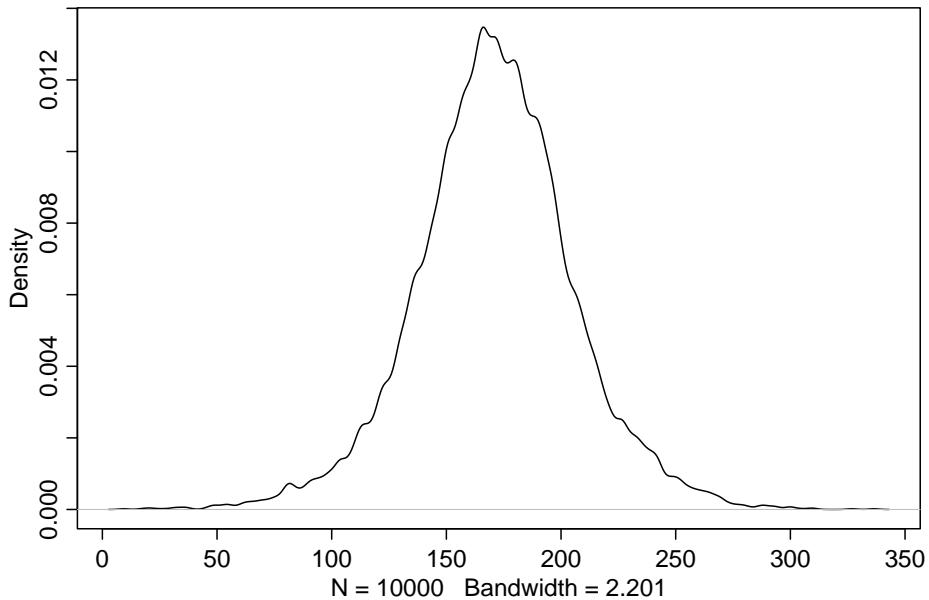
So, we simply draw μ and σ from the priors and then draw heights from the normal distribution using the drawn parameters.

Vizualisation of the prior for heights:

```
sample_mu <- rnorm(10^4, 171.1, 20)
sample_sigma <- runif(10^4, 0, 50)
prior_h <- rnorm(10^4, sample_mu, sample_sigma)
length(prior_h)
```

```
## [1] 10000
```

```
dens(prior_h)
```



The prior is not itself a Gaussian distribution, but a distribution of relative plausibilities of different heights, before seeing the data.

Note how created the model predictions for heights. We first drew μ and σ independently (there is no arrow between μ and σ) from the priors. Then we drew heights from the normal distribution using the drawn parameters. You just follow the model structure.

Now, there are a couple of different ways to estimate the model incorporating the new data. For didactic reasons, grid approximation is often used (as in the book). For many parameters, this approach becomes more and more infeasible (due to combinatorial explosion).

We will skip that for now and use quadratic approximation instead which works well for many common procedures in applied statistics (like linear regression). Later, you'll probably use (or the software in the background) mostly Markov chain Monte Carlo (MCMC) sampling to get the posterior. Pages 39 and the following explain the 3 concepts grid approximation, quadratic approximation and MCMC.

In short, **quadratic approximation** assumes that our posterior distribution of body heights can be approximated well by a normal distribution, at least near the peak.

Please read the addendum to get a clearer picture of what a bivariate normal distribution is.

Using the `rthinking` package we can estimate the model using quadratic approximation. First, we define the model in the `rthinking` syntax (see R code 4.25 in the book).

```
library(rethinking)
flist <- alist(
  height ~ dnorm(mu, sigma),
  mu ~ dnorm(171.1, 20),
  sigma ~ dunif(0, 50)
)
```

Then we estimate/fit the model using quadratic approximation.

```
m_heights <- quap(flist, data = d2)
```

Now let's take a look at the fitted model: (Note: In the online-version of the book, they used the command `map` instead of `quap`.)

The `precis` function displays concise parameter estimate information (from the posterior) for an existing model fit.

```
precis(m_heights)
```

```
##           mean        sd      5.5%     94.5%
## mu    154.603808 0.4119493 153.945433 155.262182
## sigma 7.730487 0.2913063  7.264924  8.196051
```

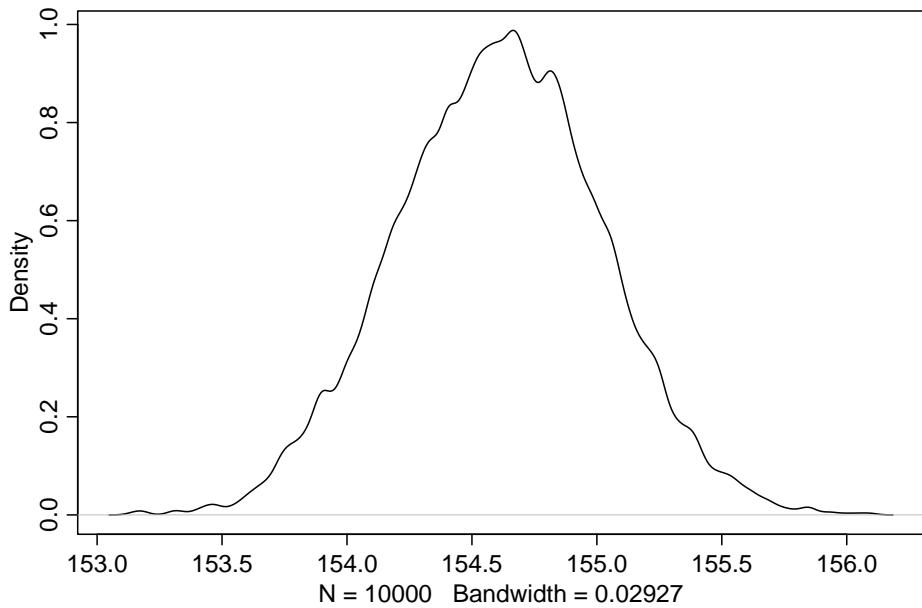
Above, we see the mean of the posterior for μ and σ ; and a **89% credible interval** for those parameters. Note that these are rather tight credible intervals. We are rather confident that the mean is somewhere between 154 and 155 cm and the standard deviation is between 7 and 8 cm.

We can now plot the posterior distribution of the mean (μ) and the standard deviation (σ) separately by drawing from the posterior distribution.

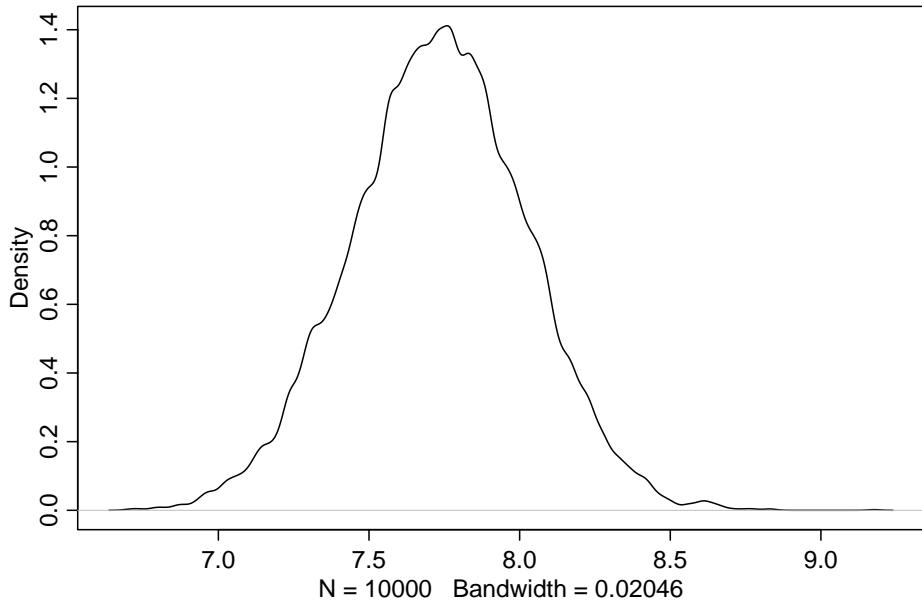
```
post <- extract.samples(m_heights, n = 10^4)
head(post)
```

```
##           mu      sigma
## 1 154.4298 7.746984
## 2 154.6450 7.533921
## 3 154.7493 7.798503
## 4 154.8317 7.477707
## 5 154.8705 7.979073
## 6 154.1648 7.932372
```

```
dens(post$mu)
```



```
dens(post$sigma)
```



Note, that **these samples come from a multi-dimensional posterior distribution**. In our case, we approximated the **joint** posterior distribution of μ

and σ with a bivariate normal distribution. They are not necessarily independent from each other, but in this case they are (see exercise 6). We know this from the prior definition above. μ and σ are both defined as normal respectively uniform distributions and by definition do not influence each other. This is also visible in the vizualisation of the model structure: There is no confounding variable or connection between those priors. One could think of a common variable Z that influences both μ and σ . This could be genetic similarity which could influence both μ and σ .

Let's verify that μ and σ are uncorrelated:

```
vcov(m_heights)
```

```
##           mu      sigma
## mu    0.1697022056 0.0001473602
## sigma 0.0001473602 0.0848593685
```

gives you the variance-covariance matrix of the parameters of the posterior distribution. In the diagonal you see the variance of the parameters.

```
diag(vcov(m_heights))
```

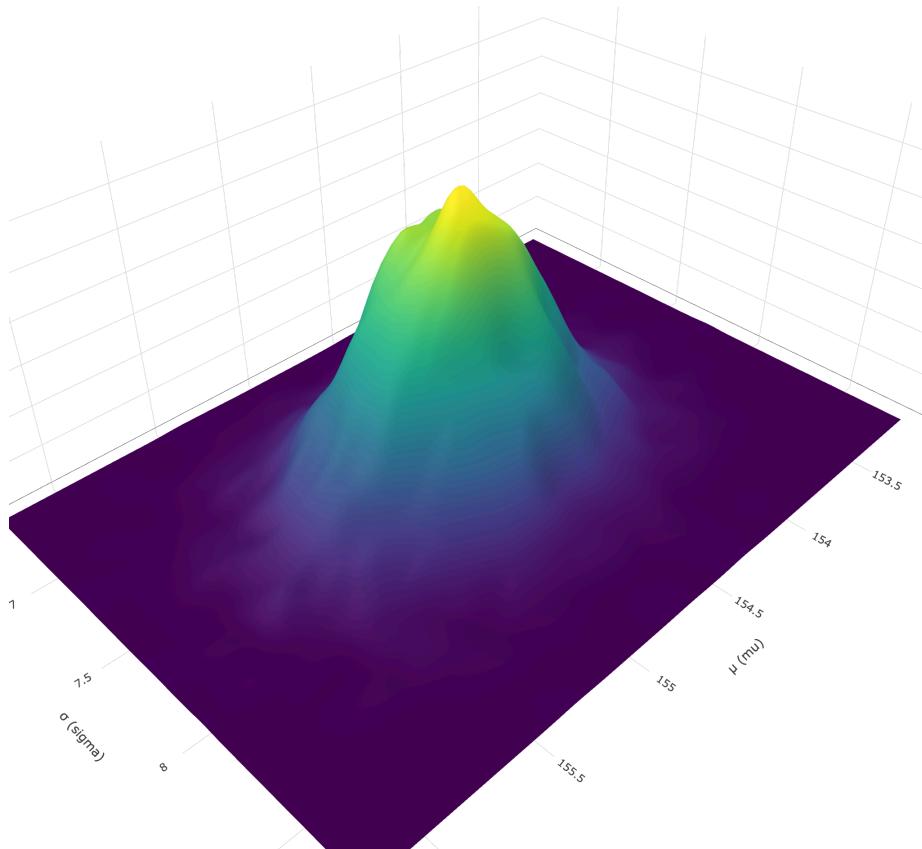
```
##           mu      sigma
## 0.16970221 0.08485937
```

And we can compute the correlation matrix easily:

```
cov2cor(vcov(m_heights))
```

```
##           mu      sigma
## mu    1.000000000 0.001227967
## sigma 0.001227967 1.000000000
```

Let's plot the posterior in 3D, because we **can**:

3D Posterior Density of μ and σ 

How beautiful ist that?

This shows how credible each combination of μ and σ is based on our priors and the data observed. The higher the mountain for a certain parameter combination, the more credible this combination is.

We see in the 3D plot, that the “mountain” is not rotated, indicating graphically that the parameters are independent from each other.

We also see in the correlation matrix, the correlation of the parameters is ~ 0 . In the context of a joint normal distribution, this means that the parameters are also independent.

And, it is not an accident that the posterior looks like this. Using quadratic

approximation, we used the bivariate normal distribution to **approximate** the posterior.

2.3 Classical approach for the simplest model

We have seen, how we could use prior knowledge to fit a very simple model for body heights of a population (!Kung San) in the Bayesian framework.

Now, let's start at the same point in the classical framework. Here, we do not use any prior knowledge, at least not that explicitly.

The classical approach to fit a regression line is the so-called **least squares method**.

There are hundreds of videos online explaining this method in great detail with animations. Maybe watch these videos later, when we add a predictor to the mean model, since most of instructional videos start at the simple linear regression using two parameters (intercept (β_0 or α) and slope (β_1)).

The (**simple mean-**) **model** is:

$$Y_i = \text{height}_i = c + \varepsilon_i$$

- for some $c \in \mathbb{R}$ and
- normally distributed errors $\varepsilon_i \sim \text{Normal}(0, \sigma)$.

The errors ε_i are on average zero and have a constant standard deviation of σ . So, we assume there is a fixed, but unknown, constant c that we want to estimate and we assume that there is a special sort of error in our model that is normally distributed. Sometimes there is a large deviation from the true c , sometimes there is a small deviation. On average, the deviations are zero and the **errors should also be independent from each other**:

$$\varepsilon_i \perp \varepsilon_j \text{ for } i \neq j$$

This means that just because I have just observed a large deviation from the true c does not mean, that the probability of a large deviation in the next observation is higher/lower. Note, that we cannot readily define different types of errors in the classical framework.

But what is c ? We determine the shape of the model ourselves (constant model, or mean model) and then estimate the parameter c . By defining the shape of the model ourselves and imposing a distribution where we want to estimate the parameter of said distribution, we are in **parametric statistics**.

We choose the c which minimizes the sum of squared errors from the actual heights. This has the advantage that deviations upper and lower from the

actual height are equally weighted. The larger the deviation the (quadratically) larger the penalty.

Why do we do that? Because, if the model assumptions (more on that later) are correct, the least squares estimator is a really good estimator. How good? Later...

We want to minimize the following function:

$$\begin{aligned} SSE \text{ (Sum of Squared Errors)} (c) &= (height_1 - c)^2 + (height_2 - c)^2 + \dots + (height_n - c)^2 = \\ &= \sum_{i=1}^n (height_i - c)^2 \end{aligned}$$

The SSE is a function of c and we want to find the c that minimizes the function. Since it is a quadratic function, we can always find the minimum. We have learnt in school how to do this (hopefully): Take the derivative of the function and set it to zero. Solve for c and you have the c which yields the minimum of SSE(c).

Let's do that:

$$\begin{aligned} \frac{d}{dc} SSE(c) &= 2(height_1 - c)(-1) + 2(height_2 - c)(-1) + \dots + 2(height_n - c)(-1) = \\ &= -2 \sum_{i=1}^n (height_i - c) \end{aligned}$$

This should be zero for the minimum:

$$\begin{aligned} -2 \sum_{i=1}^n (height_i - c) &= 0 \\ \sum_{i=1}^n (height_i - c) &= 0 \\ \sum_{i=1}^n height_i - n \cdot c &= 0 \\ \hat{c} &= \frac{1}{n} \sum_{i=1}^n height_i = \overline{height_i} \end{aligned}$$

The hat over the c indicates that this is the estimated value of the true but unknown c . Everytime we estimate a parameter, we put a hat over it.

And voilà, we have estimated the parameter c of the model, which is just the sample mean of all the heights. In contrast to before, we did not put in a lot of prior knowledge, but just estimated the parameter from the data.

In R, we can do this easily:

```

mod <- lm(height ~ 1, data = d2)
summary(mod)

##
## Call:
## lm(formula = height ~ 1, data = d2)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -18.0721  -6.0071  -0.2921   6.0579  24.4729
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)
## (Intercept) 154.5971     0.4127  374.6    <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 7.742 on 351 degrees of freedom

dim(d2)

## [1] 352   4

mean(d2$height) # same as the intercept

## [1] 154.5971

sd(d2$height) / sqrt(nrow(d2)) # standard error of the estimator

## [1] 0.4126677

# test-statistic for the intercept:
mean(d2$height) / (sd(d2$height) / sqrt(nrow(d2)))

## [1] 374.6285

# residual standard error:
sqrt(sum(mod$residuals^2) / (nrow(d2) - 1))

## [1] 7.742332

```

The `~1` means that there is just a so-called **intercept** in the model. There are **no covariates**, just the constant c . This is the simplest we can do. `lm` stands for linear model and with this base command in R we ask the software to do the least squares estimation for us.

Let's look at the **R-output** of the model estimation:

- `lm(formula = height ~ 1, data = d2)`: This is the model we estimated.
- **Residuals**: The difference between the actual height and the estimated height: $r_i = \text{height}_i - \hat{c}$. A univariate 5-point summary is given.
- **Coefficients**: The estimated coefficients of the model. In this case, there is just the intercept. We get the
 - **Std. Error** of the estimate, i.e. the standard error (SE) of the mean, which is (according to the Central Limit Theorem)

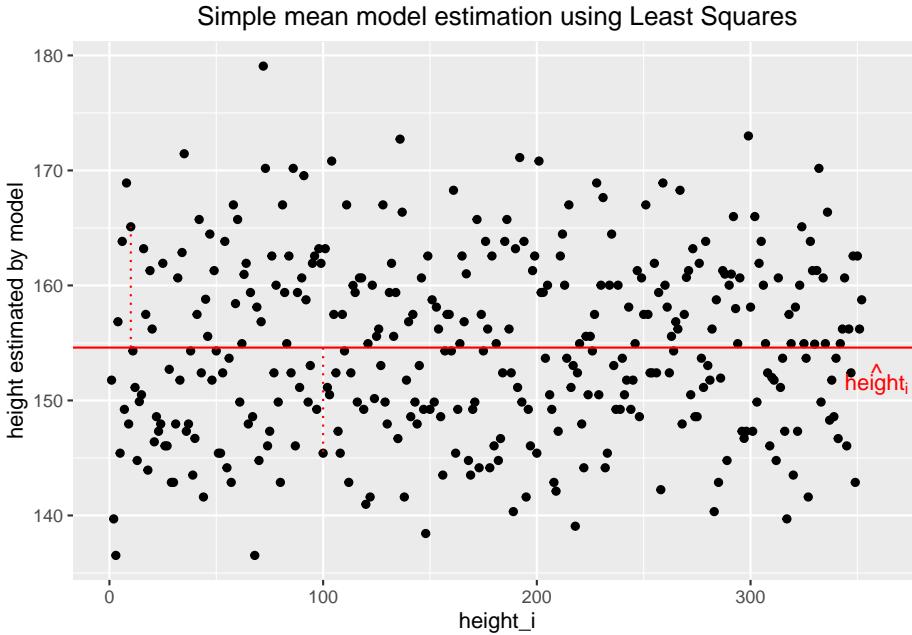
$$\frac{\sigma}{\sqrt{n}}$$

and can be estimated by the sample standard deviation divided by the square root of the sample size. $\hat{SE} = \frac{s}{\sqrt{n}}$

- the **t value** and the $\Pr(>|t|)$ which is the p -value of the (Wald-)test of the null hypothesis that the coefficient is zero ($H_0 : \text{intercept} = 0$). This is a perfect example of an absolutely useless t -test. Why? Because obviously (exercise 2) the population mean of body heights is not zero.

- **Residual standard error**: The standard deviation of the residuals $r_i = \text{height}_i - \hat{c}$. In this case identical with the sample standard deviation of heights (exercise 3). 351 degrees of freedom. There are 352 observations and 1 parameter estimated (intercept/mean). Hence, there are $352 - 1 = 351$ freely movable variables in the statistic of the sample standard deviation.

Let's look at the situation graphically:



Above, the heights are plotted against the index of the observation (the order does not matter). The variability of heights around the regression line (constant in this case) seems to stay constant, which is a good sign. We will call this **homoscedasticity** later. The dashed vertical red lines show two residuals (one > 0 , the other < 0), the difference between the actual height and the estimated height. The model-estimated heights ($\hat{heights}_i$) are all identical and nothing but the mean of all heights.

Peter Westfall explains in his excellent book a **conditional distribution approach** to regression, which is just what happens in the classical linear regression model. I **highly recommend** reading the first chapters.

What does this mean in this context? This means, that for every fixed value of the predictor (which we formally do not have), the distribution of the response is normal with mean \hat{c} and standard deviation σ . Since, we do not have a predictor (apart from the intercept), we can say, that the distribution of the heights is normal with mean \hat{c} and standard deviation σ . This is what we assumed in the model. It can also directly seen in the formula:

$$height_i = c + \varepsilon_i$$

If you add a normally distributed random variable (ε_i) to a constant (c), the result is a normally distributed. No surprise here.

We can also create new data (heights) from this model and compare the distributions of the actual data versus the model simulated data, since we have estimated the parameter c and the standard deviation σ .

```

c_hat <- mean(d2$height)
sigma_hat <- sqrt(sum(mod$residuals^2) / (nrow(d2) - 1))

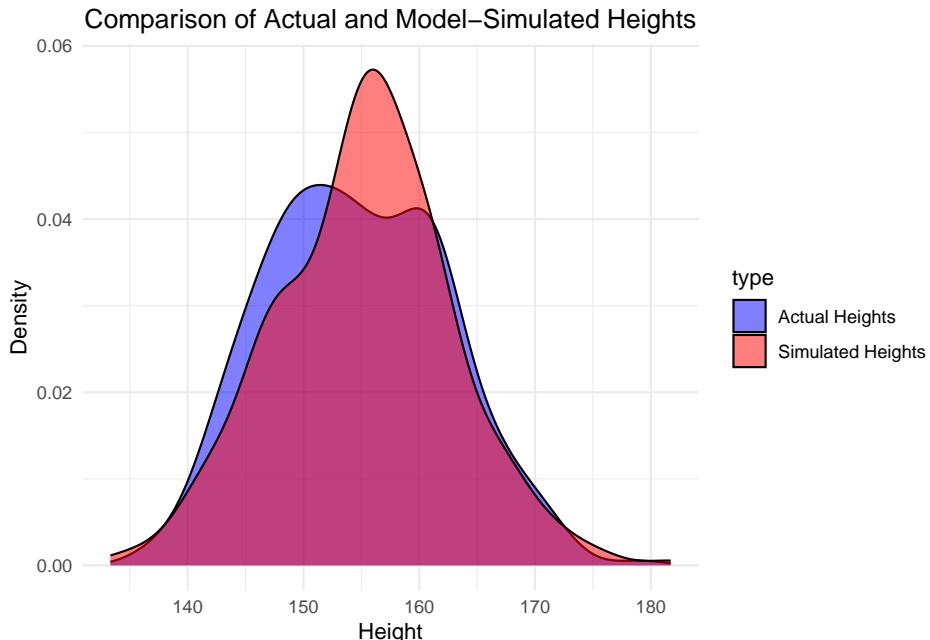
# simulate heights from model
heights_sim <- rnorm(nrow(d2), c_hat, sigma_hat) # as many as in orig.

# Convert to data frames for plotting
df_actual <- data.frame(height = d2$height, type = "Actual Heights")
df_simulated <- data.frame(height = heights_sim, type = "Simulated Heights")

# Combine the datasets
df_combined <- rbind(df_actual, df_simulated)

ggplot(df_combined, aes(x = height, fill = type)) +
  geom_density(alpha = 0.5) +
  labs(title = "Comparison of Actual and Model-Simulated Heights",
       x = "Height",
       y = "Density") +
  scale_fill_manual(values = c("blue", "red")) +
  theme_minimal() +
  theme(plot.title = element_text(hjust = 0.5))

```



One can repeat the simulation and see how the distributions changes to get a feeling for the variability of the model/data.

2.4 Exercises

[E] Easy, [M] Medium, [H] Hard

(Some) solutions to exercises can be found in the git-repo here.

2.4.1 [E] Exercise 1

Use the Swiss body heights data to determine

- the 95% “Vertrauensintervall” for μ and
- calculate the standard deviation of the heights from 21,873 Swiss people.
- Read the definition of the confidence interval in the footer of the table and explain why this is correct.

2.4.2 [E] Exercise 2

Why do we **not** need a hypothesis test to know that the population mean of body heights is not zero? Give 2 reasons.

2.4.3 [M] Exercise 3

Verify analytically that the **Residual standard error** is identical with the sample standard deviation of the heights.

2.4.4 [M] Exercise 4

Repeat the Bayesian and Frequentist estimation of the simple model using a different data set about chicken weights, which is included in R.

- Set useful priors for the mean and standard deviation of the model for the Bayesian and the frequentist version considering your a priori knowledge about chicken weights.

2.4.5 [M] Exercise 5

Try to verify the simulation of probability in the addendum below.

2.4.6 [M] Exercise 6

Independence of two random variables X and Y implies that the correlation between them is zero. A correlation of zero does not necessarily imply independence. Verify this counterexample:

For example, suppose the random variable X is symmetrically distributed about zero, and $Y = X^2$. Then Y is completely determined by X , so that X and Y are perfectly dependent, but their correlation is zero.

2.5 Addendum

2.5.1 The bivariate normal distribution

As a refresher, you can look into the old QM1 script and read the chapter “4.7 Gemeinsame Verteilungen”. Maybe this video also helps.

The bivariate normal distribution is a generalization of the normal distribution to two dimensions. Now, we look at the distribution of two random variables X and Y **at the same time**.

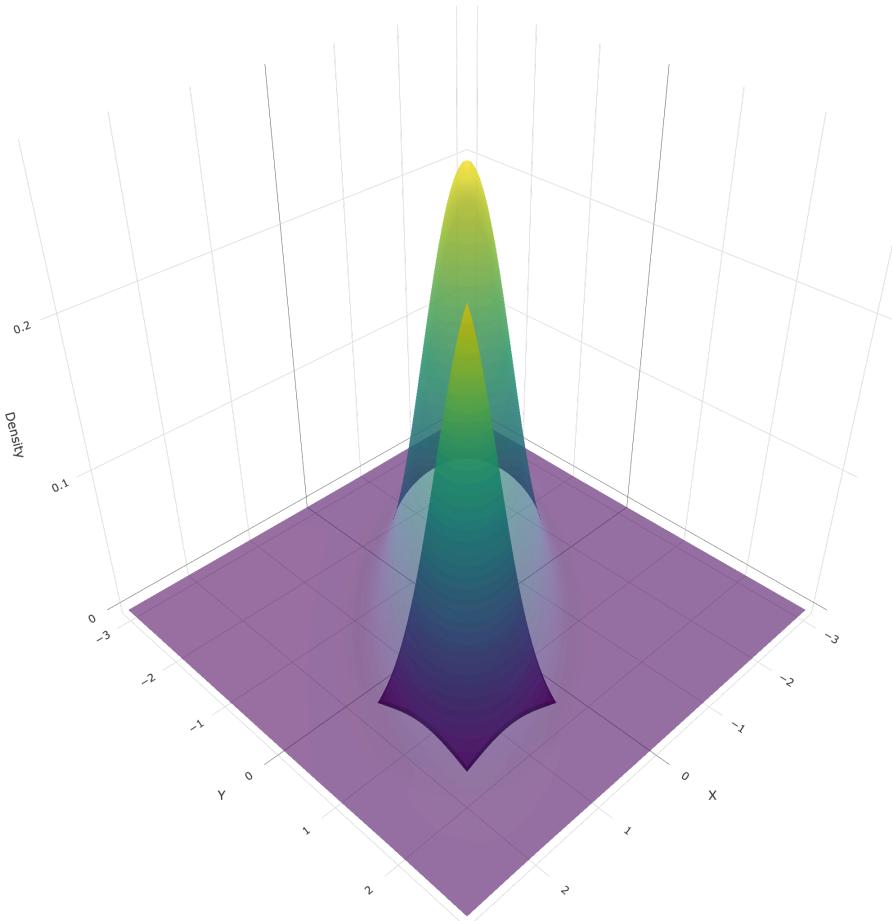
Instead of one Gaussian curve, we have a 3D curve. This curve defines how plausible different combinations of X and Y are.

Single points (like (3,6)) still have probability zero, because now the **volume** over a single point (x, y) is zero. The probability of a certain area is now the **volume** under the curve compared to the **area** under the density curve in the one-dimensional case.

Example: The following plot shows the density of a bivariate normal distribution of two variables X and Y with $\mu_X = 0$, $\mu_Y = 0$, $\sigma_X = 1$, $\sigma_Y = 1$ and $\rho = \frac{2}{3}$.

Below is the correlation matrix of the bivariate normal distribution.

```
##          [,1]      [,2]
## [1,] 1.0000000 0.6666667
## [2,] 0.6666667 1.0000000
```



If you move the plot around with your mouse, you see that there is a positive correlation between X and Y ($\rho = \frac{2}{3}$). This means that if X is above its mean, Y is also more likely to be above its mean. The variances of X and Y are both 1. That means, that if you cut through the plot in $X = 0$ or $Y = 0$, you see the same form of normal distribution. If you look at it from above, we have highlighted the section on the surface over the area $X \in [0.5, 2]$ and $Y \in [0.5, 2]$. The volume over this area under the density curve is the probability of this area: $P(X \in [0.5, 2] \text{ and } Y \in [0.5, 2])$

Calculate with R this probability with R:

```
# Load the mvtnorm package
library(mvtnorm)
```

```

# Define the parameters of the bivariate normal distribution
mu <- c(0, 0) # Mean
sigma <- matrix(c(0.75, 0.5, 0.5, 0.75), ncol = 2) # Covariance matrix

# Define the bounds of the square
highlight_x <- c(0.5, 2)
highlight_y <- c(0.5, 2)
# Calculate the probability using pmvnorm
pmvnorm(
  lower = c(highlight_x[1], highlight_y[1]),
  upper = c(highlight_x[2], highlight_y[2]),
  mean = mu,
  sigma = sigma
)

## [1] 0.1526031
## attr(,"error")
## [1] 1e-15
## attr(,"msg")
## [1] "Normal Completion"

```

Since we do not believe everything we are told, we rather check via simulation, if 0.1526 is a plausible value for the probability:

```

# Load necessary library
library(MASS)

# Define the parameters of the bivariate normal distribution
mu <- c(0, 0) # Mean
sigma <- matrix(c(0.75, 0.5, 0.5, 0.75), ncol = 2) # Covariance matrix

# Define the bounds of the square
highlight_x <- c(0.5, 2)
highlight_y <- c(0.5, 2)

# Number of simulations
n_sim <- 10^4

set.seed(343434)
# Simulate bivariate normal samples
samples <- mvrnorm(n = n_sim, mu = mu, Sigma = sigma)

# Count how many samples fall within the square
inside_square <- sum(

```

```
samples[, 1] >= highlight_x[1] & samples[, 1] <= highlight_x[2] &
samples[, 2] >= highlight_y[1] & samples[, 2] <= highlight_y[2]
)

# Estimate the probability
inside_square / n_sim

## [1] 0.1557
```

Looks good.

Chapter 3

Simple Linear Regression

3.1 Simple Linear Regression in the Bayesian Framework

You can watch this video as primer.

We will now add one covariate/explanatory variable to the model. Refer to Statistical Rethinking “4.4 Linear prediction” or “4.4 Adding a predictor” as it’s called in the online version of the book.

So far, our “regression” did not do much to be honest. The mean of a list of values was already calculated in the descriptive statistics section before and we have mentioned how great this statistic is as measure of location and where its weaknesses are.

Now, we want to model **how** body height and weight are **related**. Formally, one wants to *predict* body heights from body weights.

Here and in the frequentist framework, we will see that it is **not the same** problem (and therefore results in a different statistical model) **to predict body weights from body heights or vice versa**.

We remember the following:

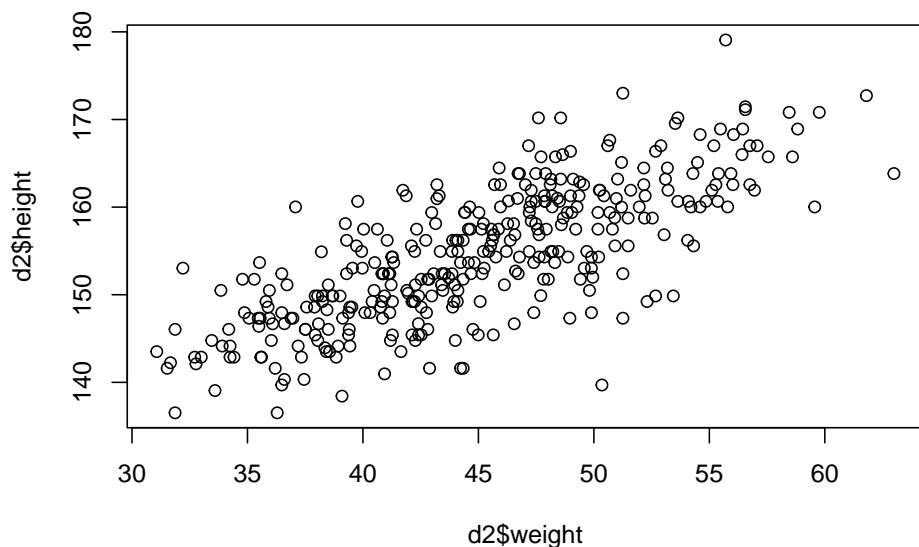
- Regress Y on X , which is equivalent to predict Y from X . We know X and want to predict Y .
- Regress X on Y , which is equivalent to predict X from Y . We know Y and want to predict X .

The word “predictor” is important here. It is a technical term and describes a variable that we know (in our case weight) and with which we want to “guess

as good as possible” the value of the dependent variable (in our case height). “As good as possible” means that we put a penalty on an error. The farer our prediction is aways from the true value (y_i), the higher the penalty. And not only that, but if you are twice as far away from the true value, you should be penalized four times as much. This is the idea behind the squared error loss function and the core of the least squares method. What if we would punish differently, you ask? There are many loss functions one could use (for instance the Huber loss), maybe we will see some later. For now, we punish quadratically.

We **always** visualize the data first to improve our understanding. First comes descriptive statistics, then one can think about modeling.

```
plot(d2$height ~ d2$weight)
```



It's not often that you see such a clean plot. The scatterplot indicates a linear relationship between the two variables. The higher the weight, the higher the height; with some deviations of course and we decide that normally distributed errors are a good idea. This relationships is neither causal, nor deterministic.

- It is not causal since an increase in weight does not necessarily lead to an increase in height, especially in grown-ups.
- It is not deterministic since there are deviations from the line. If it was deterministic, we would not need statistical modeling.

For simpler notation, we will call `d2$weight` x . \bar{x} is the mean of x .

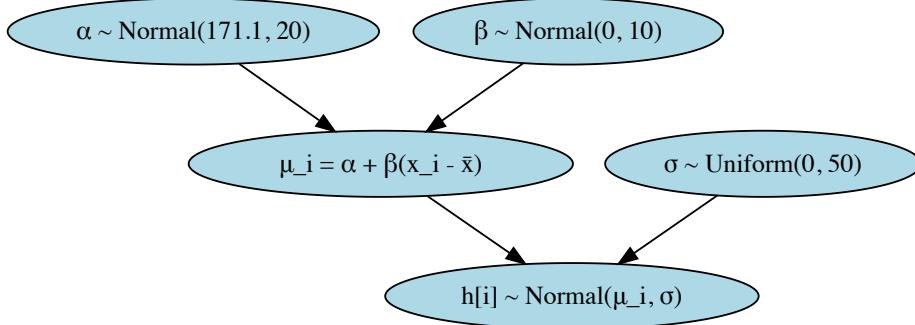
3.1.1 Model definition

Let's write down our **model** (again with the Swiss population prior mean):

$$\begin{aligned}
 h_i &\sim \text{Normal}(\mu_i, \sigma) \\
 \mu_i &\sim \alpha + \beta(x_i - \bar{x}) \\
 \alpha &\sim \text{Normal}(171.1, 20) \\
 \beta &\sim \text{Normal}(0, 10) \\
 \sigma &\sim \text{Uniform}(0, 50)
 \end{aligned}$$

Visualization of the **model structure**:

```
## file:///private/var/folders/pm/jd6n6gj10371_bml1gh8sc5w0000gn/T/Rtmpj9F7oh/file1574d52bc45be
```



There are now additional lines for the priors of α and β (compared to the simple mean model before). The model structure also shows the way to simulate from the prior. One starts at the top and ends up with the heights.

- h_i is the height of the i -th person and we assume it is normally distributed.
- μ_i is the mean of the height of the i -th person and we assume it is linearly dependent on the difference $x_i - \bar{x}$. Compared to the intercept model, a different mean is assumed for each person depending on his/her weight.
- α is the intercept and we use the same prior as before.
- β is the slope of the line and we use the normal distribution as prior for it, hence it can be positive or negative and how plausible each value is, is determined by that specific normal distribution. Note, that we could easily adapt the distribution to any distribution we like.
- The prior for σ is unchanged.
- $x_i - \bar{x}$ is the deviation of the weight from the mean weight, thereby **we center** the weight variable. This is a common practice in regression analysis. A value $x_i - \bar{x} > 0$ implies that the person is heavier than the average.

The linear model is quite popular in applied statistics and one reason is probably the rather straightforward interpretation of the coefficients: A one unit increase in weight is on average (in the mean) associated with a β unit increase/decrease (depending if β is > 0 or < 0) in height.

3.1.2 Priors

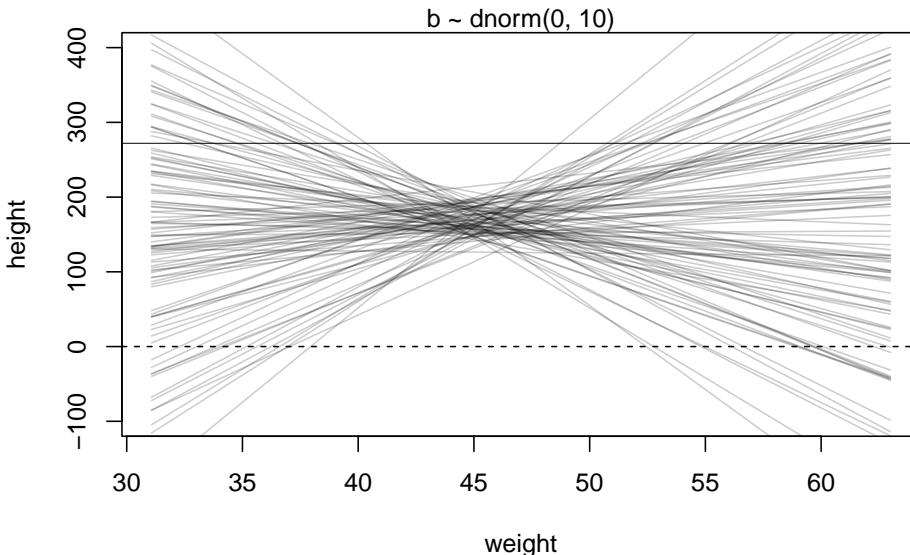
We want to plot our prior predictions to get a feeling **what the model would predict without seeing the data**. This is a kind of “sanity check” to see if the priors and the model definition are reasonable. Again, we just draw from the assumed distributions for α and β 100 times and draw the corresponding lines. Just as the model definition says.

```
set.seed(2971)
N <- 100 # 100 lines
a <- rnorm(N, 171.1, 20)
b <- rnorm(N, 0, 10)

xbar <- mean(d2$weight)

# start with empty plot
plot(NULL, xlim = range(d2$weight), ylim = c(-100, 400),
      xlab = "weight", ylab = "height")
abline(h = 0, lty = 2) # horizontal line at 0
abline(h = 272, lty = 1, lwd = 0.5) # horizontal line at 272
mtext("b ~ dnorm(0, 10)")

# Overlay the 100 lines
for (i in 1:N) {
  curve(a[i] + b[i] * (x - xbar),
         from = min(d2$weight), to = max(d2$weight),
         add = TRUE, col = col.alpha("black", 0.2))
}
```



This linear relationship defined with the chosen priors seems rather non-restrictive. According to our priors, one could see very steeply rising or falling relationships between weight and expected heights. We could at least make the priors for the slope (β) non-negative. One possibility to do this is to use a log-normal distribution for the prior of β which can only take non-negative values.

$$\beta \sim \text{Log-Normal}(0, 1)$$

Lets plot the priors again.

```
set.seed(2971)
N <- 100 # 100 lines
a <- rnorm(N, 171.1, 20)
b <- rlnorm(N, 0, 1)

xbar <- mean(d2$weight)

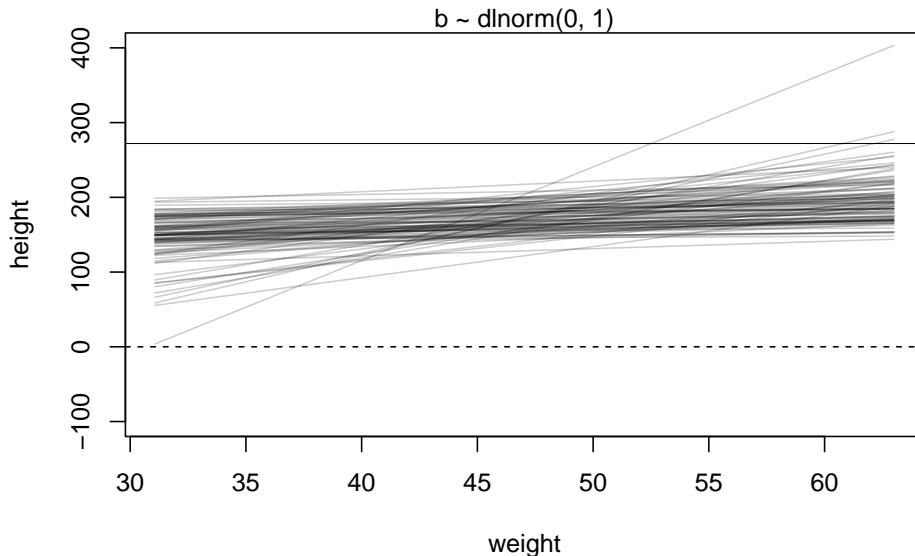
plot(NULL, xlim = range(d2$weight), ylim = c(-100, 400),
     xlab = "weight", ylab = "height")
abline(h = 0, lty = 2) # horizontal line at 0
abline(h = 272, lty = 1, lwd = 0.5) # horizontal line at 272
mtext("b ~ dlnorm(0, 1)")

# Overlay the 100 lines
for (i in 1:N) {
  curve(a[i] + b[i] * (x - xbar),
```

```

    from = min(d2$weight), to = max(d2$weight),
    add = TRUE, col = col.alpha("black", 0.2))
}

```



This seems definitely more realistic. There is some sort of positive linear relationship between weight and expected height.

3.1.3 Fit model

Now, let's estimate the posterior/fit the model as before:

```

# load data again, since it's a long way back
library(rethinking)
data(Howell1)
d <- Howell1
d2 <- d[d$age >= 18, ]
xbar <- mean(d2$weight)
# fit model
mod <- quap(
  alist(
    height ~ dnorm(mu, sigma),
    mu <- a + b * (weight - xbar),
    a ~ dnorm(171.1, 100),
    b ~ dnorm(0, 10),
    sigma ~ dunif(0, 50)
  ) ,
  data = d2)

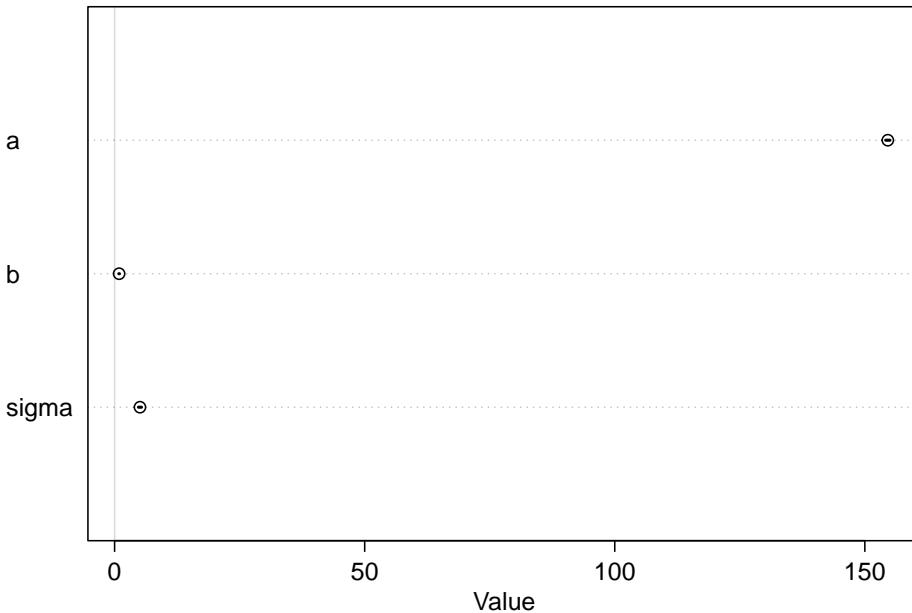
```

Note that the model definition was now directly included in the `quap` function. Let's look at the **marginal distributions** of the parameters:

```
precis(mod)
```

```
##               mean        sd      5.5%     94.5%
## a     154.5972120 0.27033045 154.1651717 155.0292523
## b      0.9050131 0.04192754  0.8380048  0.9720214
## sigma 5.0718673 0.19115323  4.7663675  5.3773671
```

```
plot(mod)
```



Note, that the credible intervals in the plot of the coefficients are hardly noticeable. The reason is that the intercept a is rather large, whereas b and σ are comparatively small. The plot is scaled to the largest parameter.

We can improve this by centering the height variable as well. Or, we could standardize both weight and height. This would change the interpretation of the β parameter. It would then be the expected change in standard deviations when changing weight by one standard deviation. See exercise 12.

The analysis yields estimates for all our parameters of the model: α , β and σ . The estimates are the mean of the posterior distribution.

See exercise 2.

Interpretation of β : The mean of the posterior distribution of β is 0.9. A person with a weight of 1 kg more weight can be expected to be 0.9 cm taller.

A 89% credible interval for this estimate is [0.83, 0.97]. We can be quite sure that the slope is positive (of course we designed it that way too via the prior).

It might also be interesting to inspect the **variance-covariance matrix**, respectively the correlation between the parameters as we did before in the intercept model. Remember, these are the correlations of parameters in the multivariate (because three parameters simultaneously) posterior distribution.

```
diag(vcov(mod))
```

```
##           a          b        sigma
## 0.073078550 0.001757918 0.036539558
```

```
round(cov2cor(vcov(mod)), 2)
```

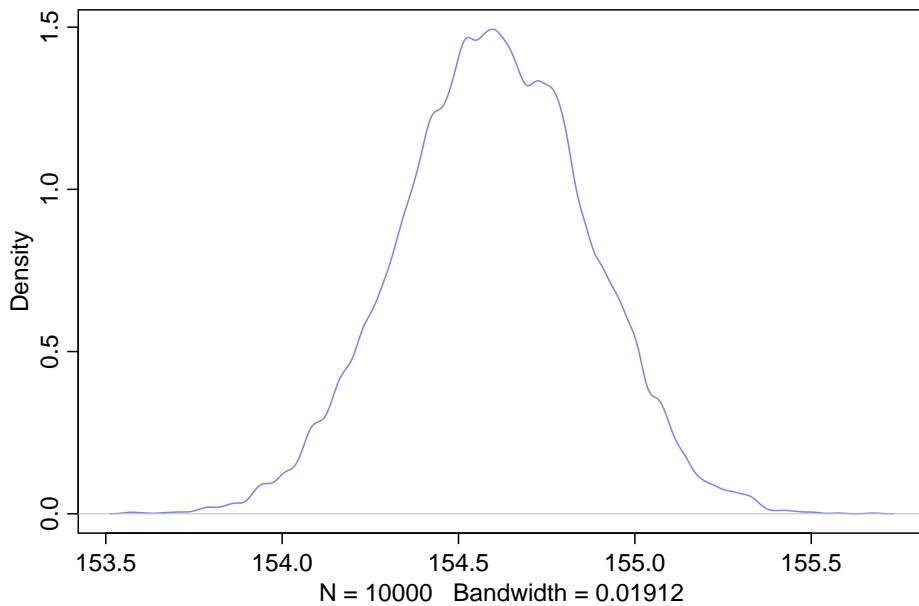
```
##      a b sigma
## a  1 0    0
## b  0 1    0
## sigma 0 0    1
```

- `diag(vcov(mod))` gives the variances of the parameters and
- `cov2cor(vcov(mod))` the correlations. As we can see the correlations are (near) zero. Compare to the graphical display of the model structure. There is no connection.

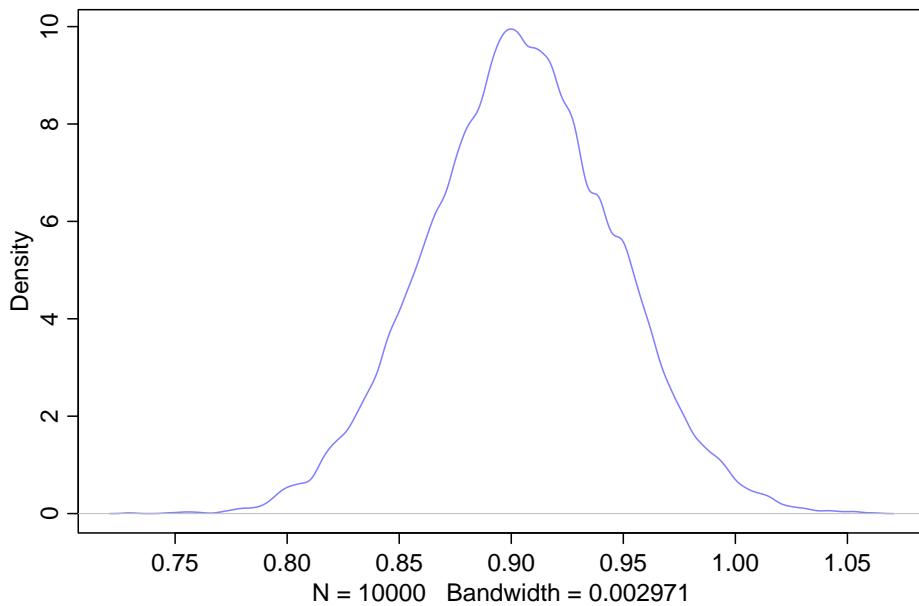
3.1.4 Result

Graphical end result of fitting the model: We plot the marginal posterior distributions of α and β , and also the raw data with the found regression line.

```
post <- extract.samples(mod)
dens(post$a, col = rangi2)
```



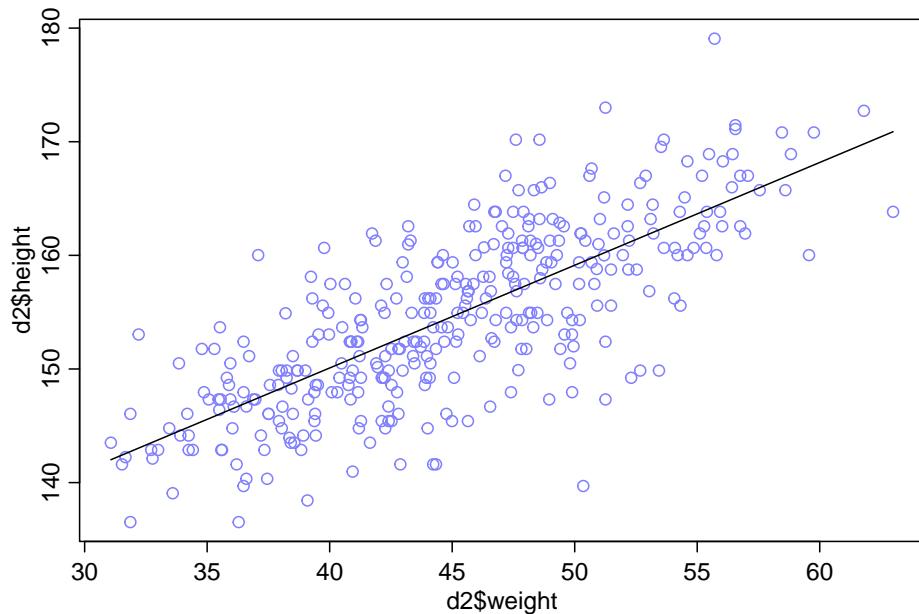
```
dens(post$b, col = rangi2)
```



```
# both posterior plots seem symmetric  
# we use the mean as point estimate  
# for a and b.  
a_quap <- mean(post$a)
```

```
b_quap <- mean(post$b)

plot(d2$height ~ d2$weight, col = rangi2)
curve(a_quap + b_quap * (x - xbar), add = TRUE)
```



3.1.5 Credible bands

We could draw again and again from the posterior distribution and calculate the means like above. Plotting the regression lines with the respective parameters α , β would indicate the variability of the estimates. Note that we do not draw from the data (as one does in bootstrap resampling), but from the posterior distribution. The `link` function does this for us. It takes the posterior distribution We have just fit, samples α and β from it, calculates the mean and then samples from this normal distribution for the mean at a given weight. Refer to pages 98-106 in the current version of the book Statistical Rethinking for all details.

```
# Define a sequence of weights for predictions
weight.seq <- seq(from = 25, to = 70, by = 1)

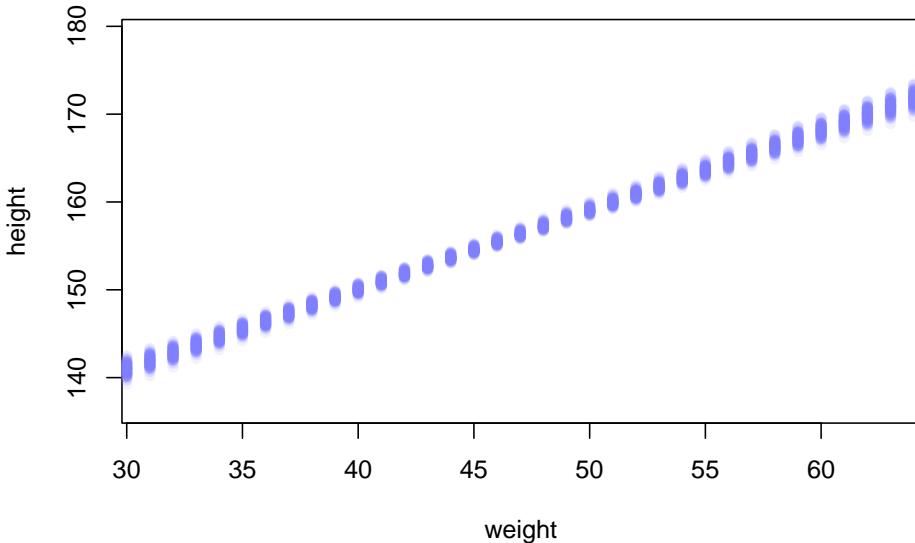
# Use the model to compute mu for each weight
mu <- link(mod, data = data.frame(weight = weight.seq))
str(mu)
```

```

##  num [1:1000, 1:46] 138 136 137 136 137 ...
## Visualize the distribution of mu values
plot(height ~ weight, d2, type = "n") # Hide raw data with type = "n"

# Loop over samples and plot each mu value
for (i in 1:100) {
  points(weight.seq, mu[i, ], pch = 16, col = col.alpha(rangi2, 0.1))
}

```



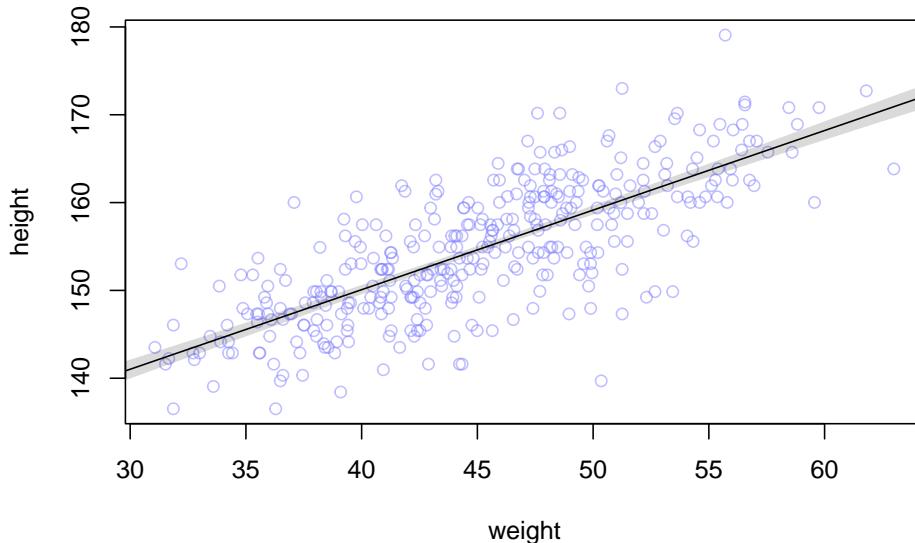
The `link` function fixes the weight at the values in `weight.seq` and draws samples from the posterior distribution of the parameters. We will do the analog thing in the frequentist framework.

We can also draw a nice shade for the regression line:

```

# Summarize the distribution of mu
mu.mean <- apply(mu, 2, mean)
mu.PI <- apply(mu, 2, PI, prob = 0.89)
plot(height ~ weight, d2, col = col.alpha(rangi2, 0.5))
lines(weight.seq, mu.mean)
shade(mu.PI, weight.seq)

```



The function `PI` from the `rethinking` package calculates the 89% percentile interval for the mean of the height at a certain weight. The `apply` function applies this function to each columns (hence the `2`) of the matrix `mu`.

As we can see, we are pretty sure about the mean of height which we wanted to model in the first place.

Mean modeling is one thing, individual prediction is another. Given a certain weight of a person, what is the height of the same person? The first line in the model definition ($height_i \sim Normal(\mu_i, \sigma)$) tells us that a person's weight is distributed *around* the mean (which linearly depends on weight) and is not necessary the mean itself.

To get to an **individual prediction**, we need to consider the uncertainty of the parameter estimation *and* the uncertainty from the Gaussian distribution around the mean (at a certain weight). We do this with `sim`.

```
# Simulate heights from the posterior
sim.height <- sim(mod, data = list(weight = weight.seq))
str(sim.height)
```

```
## num [1:1000, 1:46] 138 130 130 147 136 ...
```

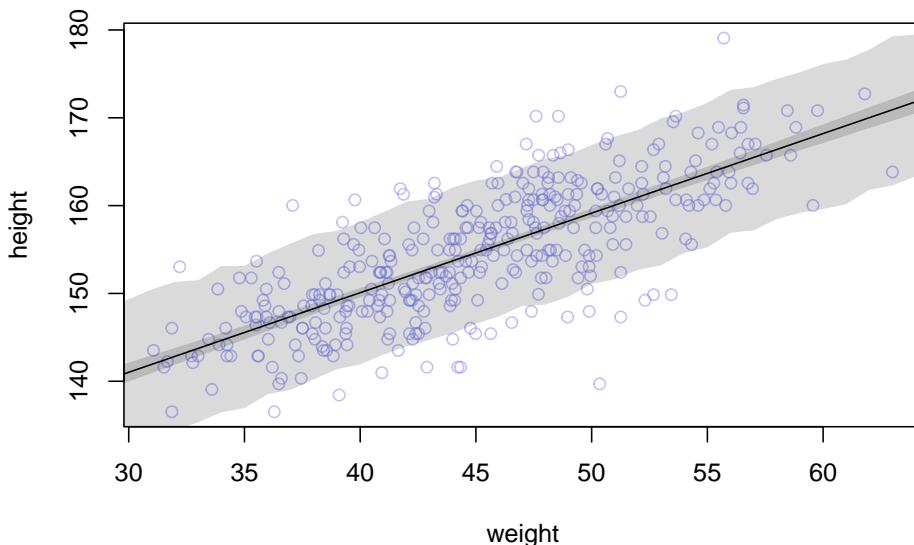
```
# Compute the 89% prediction interval for simulated heights
height.PI <- apply(sim.height, 2, PI, prob = 0.89)

# Plot the raw data
plot(height ~ weight, d2, col = col.alpha(rangi2, 0.5))
```

```
# Draw MAP (mean a posteriori) line
lines(weight.seq, mu.mean)

# Draw HPDI (highest posterior density interval) region for mu
shade(mu.PI, weight.seq)

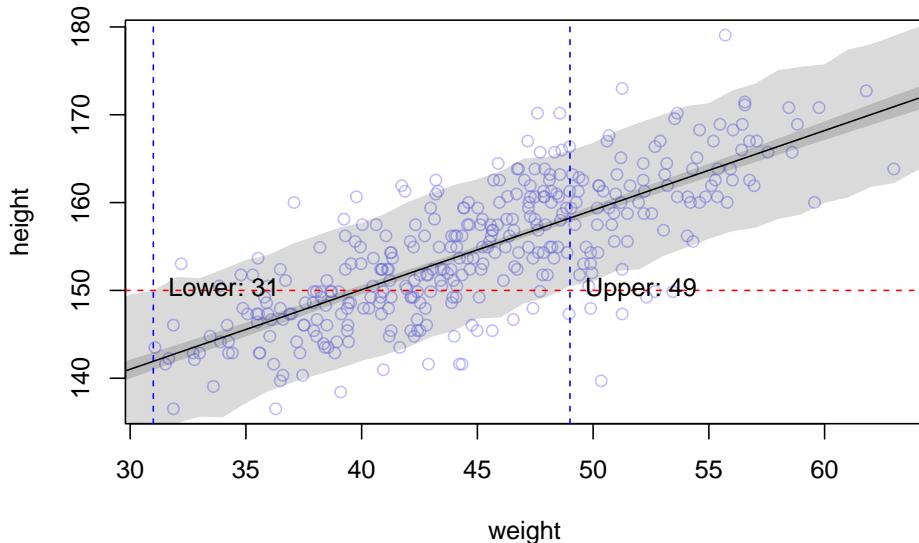
# Draw PI (prediction interval) region for simulated heights
shade(height.PI, weight.seq)
```



Here, the `PI` function is applied to the simulated heights and calculates the 89% percentile interval for each weight.

The lighter and wider shaded region is where the model expects to find 89% of the heights of a person with a certain weight.

This part is **sometimes a bit disillusioning** when seen for the first time: Draw a horizontal line at 150 cm and see how many weights (according to the individual prediction) are compatible with this height. Weights from 30 to 50 kg are compatible with this height according to the 89% prediction interval:



The higher the credibility, the wider the interval, the wider the range of compatible weights. In our example, more than 60% of the weight-range are plausible to predict a height of 150 cm.

```
(50 - 30) / (range(d2$weight)[2] - range(d2$weight)[1])
```

```
## [1] 0.6265362
```

On the other hand: We did not model the relationship this way. We modeled **height depending on weight** and not the other way around. In the next chapter, we will regress weight on height (yes, this is the correct order) and see what changes.

3.1.6 Summary

- We have added a covariate (weight) to the simple mean model to predict height.
- We have centered the weight variable.
- We have defined and refined priors for the intercept and slope.
- We have estimated the posterior distribution of the parameters using quadratic approximation with `quap`.
- We have visualized the result.
- We have created credible bands for mean and individual predictions.

3.2 Simple Linear Regression in the Frequentist Framework

We will now do the same analysis in the frequentist framework while introducing some foundational theory along the way. I recommend reading the first couple of chapters from Westfall.

3.2.1 Model definition

Our linear model is defined as:

$$h_i = \beta_0 + \beta_1 x_i + \varepsilon_i$$

where

- ε_i is the error term with $\varepsilon_i \sim N(0, \sigma^2), \forall i$
- β_0 is the unknown but fixed intercept
- β_1 is the unknown but fixed slope

3.2.1.1 Model Assumptions of the Classical Regression Model (Westfall, 1.7):

The first and **most important assumption** is that the data are produced probabilistically, which is specifically stated as

$$Y|X = x \sim p(y|x)$$

What does this mean?

- $Y|X = x$ is the random variable Y **conditional** on X being equal to x, i.e. the distribution of Y if we know the value of X (in our example the weight in kg). This is a nice image of what is meant here.
- $p(y|x)$ is the distribution of potentially observable Y given $X = x$. In our case above this was the normal distribution with mean μ_i and variance σ .

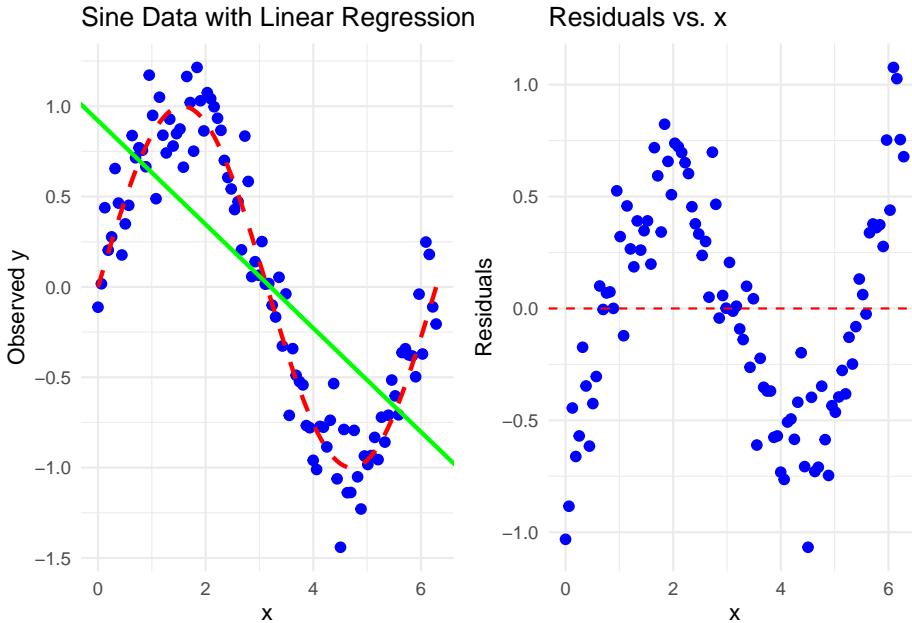
You can play with this shiny app to improve your understanding. It offers the option “Bedingte Verteilung anzeigen”.

One always thinks about the so-called data generating process (Westfall, 1.2). How did the data come about? There is a process behind it and this process is attempted to be modeled.

Further assumptions:

- **Correct functional specification:** The conditional mean function $f(x) = \mathbb{E}(Y|X = x)$. In the case of the linear model, the assumption is $\mathbb{E}(Y|X = x) = \alpha + \beta x$. The **expectation** of Y (height) depends linearly on x (weight). This assumption is violated when the true relationship is not linear or the data at least suggest that it is not linear, like here.
- **The errors are homoscedastic** (constant variance σ^2). This means the variances of all conditional distributions $p(y|x)$ are constant ($= \sigma^2$). This assumption is (for instance) violated if points are spreading out more and more around the regression line, indicating that the errors are getting larger.
- **Normality.** For the classical linear regression model all the conditional distributions $p(y|x)$ are normal distributions. It could well be, that the errors are not nicely normally distributed around the regression line, for instance if we have a lot of outliers upwards and the distribution is skewed, like here (Figure 5.20).
- The **errors are independent** of each other. The potentially observable $\varepsilon_i = Y_i - f(\mathbf{x}_i,)$ is uncorrelated with $\varepsilon_j = Y_j - f(\mathbf{x}_j,)$ for $i \neq j$. This assumption is violated if the errors are correlated, here is an example: The true data comes from a sine curve and we estimate a linear model (green), which does not fit the data well (left plot). The residuals plot shows clear patterns (right plot) and indicates that the errors are correlated. Specifically, the errors around $x = 2$ and $x = 4$ are negatively correlated (see exercise 6).

```
##  
## Attaching package: 'patchwork'  
  
## The following object is masked from 'package:MASS':  
##  
##      area
```



In the case above, the errors are **not conditionally independent**. If we condition on $X = 2$ and $X = 4.5$, the errors are correlated ($r \sim -0.3$), which they should not be.

These assumptions become clearer as we go along and should be checked for every model we fit. They are not connected, they can all be true or false. The question is not “Are the assumptions met?” since they never are exactly met. The question is **how “badly”** the assumptions are violated?

Remember, **all models are wrong, but some are useful**.

In full, the classical linear regression model can be written as:

$$Y_i | X_i = x_i \sim_{independent} N(\beta_0 + \beta_1 x_{i1} + \dots \beta_k x_{ik}, \sigma^2)$$

for $i = 1, \dots, n$.

3.2.2 Fit the model

Again, we fit the model using the least squares method. For a neat animated explanation, visit this video. There are literally hundreds of videos on the topic. Choose wisely. Not all are good. If in doubt, use our recommended books as reading materials. This is the most reliable source. A hint along the way: Be very sceptical if you ask GPT about information, although for this special case one has a good chance of getting a decent answer due to the vast amount of training data.

One has to minimize the sum of squared differences between the true heights and the model-predicted heights in order to find β_0 and β_1 .

$$SSE(\beta_0, \beta_1) = \sum_{i=1}^n (y_i - (\beta_0 + \beta_1 x_i))^2$$

We omit the technical details (set derivative to zero and solve the system) and give the results for β_0 and β_1 :

$$\begin{aligned}\hat{\beta}_0 &= \bar{y} - (\hat{\beta}_1 \bar{x}), \\ \hat{\beta}_1 &= \frac{\sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y})}{\sum_{i=1}^n (x_i - \bar{x})^2} = \frac{s_{x,y}}{s_x^2} = r_{xy} \frac{s_y}{s_x}.\end{aligned}$$

where:

- r_{xy} is the sample correlation coefficient between x and y
- s_x and s_y are the uncorrected sample standard deviations of x and y
- s_x^2 and s_{xy} are the sample variance and sample covariance, respectively

Interpretation of $\hat{\beta}_0$ and $\hat{\beta}_1$: see exercise 3.

Let's use R again to solve the problem:

```
library(rethinking)
data(Howell1)
d <- Howell1
d2 <- d[d$age >= 18, ]
mod <- lm(height ~ weight, data = d2)
summary(mod)

##
## Call:
## lm(formula = height ~ weight, data = d2)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -19.7464  -2.8835   0.0222   3.1424  14.7744
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)    
## (Intercept) 113.87939   1.91107  59.59 <2e-16 ***
## weight       0.90503    0.04205  21.52 <2e-16 ***
## ---
```

```

## Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 5.086 on 350 degrees of freedom
## Multiple R-squared: 0.5696, Adjusted R-squared: 0.5684
## F-statistic: 463.3 on 1 and 350 DF, p-value: < 2.2e-16

```

Interpretation of R-output:

- **Call:** The model that was fitted.
- **Residuals:** $r_i = \text{height}_i - \widehat{\text{height}}_i$. Differences between true heights and model-predicted heights.
- **Coefficients:** Estimated for β_0 and β_1 . We call them $\hat{\beta}_0$ and $\hat{\beta}_1$.
 - **Estimate:** The (least squares) estimated value of the coefficient.
 - **Std. Error:** The standard error of the estimate.
 - **t value:** The value of the t -statistic for the (Wald-) hypothesis test $H_0 : \beta_i = 0$.
 - **Pr(>|t|):** The p -value of the hypothesis test.
- **Residual standard error:** The estimate of σ which is also a model parameter (as in the Bayesian framework).
- **Multiple R-squared:** The proportion of the variance explained by the model (we will explain this below).
- **Adjusted R-squared:** A corrected version of the R^2 which takes into account the number of predictors in the model.
- **F-statistic:** The value of the F -statistic for the hypothesis test: $H_0 : \beta_1 = \beta_2 = \dots = \beta_k = 0$. Note, the alternative hypotheses to this test is that *any* of the β_i is not zero. If that is the case, the model explains more than the mean model with just β_0 .

We could also **solve the least squares problem graphically**: We want to find the values of β_0 and β_1 that minimize the sum of squared differences which can be plotted as 3D function. Since the function is a sum of squared terms, we should expect a paraboloid form. All we have to do is to ask R which of the coordinates (β_0, β_1) minimizes the sum of squared errors. The result confirms the results from the `lm` function. The dot in red marks the spot (Code is in the git repository):

—COMPILE CODE AT DEPLOYMENT (takes a bit)—

3.2.3 Confidence Intervals of coefficients (frequentist)

You can get CI's conveniently with the `confint` function:

```
confint(mod, level = 0.96)

##               2 %      98 %
## (Intercept) 109.939864 117.8189232
## weight       0.818351  0.9917072
```

Remember, these are frequentist confidence intervals.

If one repeats the experiment many times, the true but unknown value of the parameter will be in the interval in 96% of the cases

We can also use the simple bootstrap. The advantage of this technique is that we can basically always use it, no matter how complicated the estimator is. We do not need formulae. We simply

- create 1000 bootstrap samples,
- fit the model,
- store the coefficients.
- The 2% and 98% quantiles of the coefficients constitute the 96% bootstrap confidence interval.

```
set.seed(123)
n <- nrow(d2)
B <- 1000
boot_coefs <- matrix(NA, nrow = B, ncol = 2)
for (i in 1:B) {
  boot_idx <- sample(1:n, replace = TRUE)
  boot_mod <- lm(height ~ weight, data = d2[boot_idx, ])
  boot_coefs[i, ] <- coef(boot_mod)
}
#head(boot_coefs)
t(apply(boot_coefs, 2, quantile, c(0.02, 0.98)))

##               2%      98%
## [1,] 110.1862455 117.4516455
## [2,]  0.8229982  0.9859997
```

The CIs are quite similar to the ones from the `confint` function.

In the Bayesian setting, we used the centered weight variable. Let's do this here too for comparison and use 89% coverage probability.

```
d2$weight_centered <- d2$weight - mean(d2$weight)
mod_centered <- lm(height ~ weight_centered, data = d2)
#summary(mod_centered)
confint(mod_centered, level = 0.89)
```

```
##           5.5 %      94.5 %
## (Intercept) 154.162715 155.0314698
## weight_centered 0.837658  0.9724002
```

Compare with `precis` from the Bayesian model:

```
##          mean        sd      5.5%     94.5%
## a    154.5972131 0.27033041 154.165173 155.0292533
## b     0.9050133 0.04192753  0.838005  0.9720216
## sigma 5.0718667 0.19115317  4.766367  5.3773663
```

We are glad to see that both analyses align really nicely.

3.2.4 ANOVA (Analysis of Variance)

A non-obvious and very useful finding is that the total variability (SST) in the data (our heights) can be **decomposed** (or analysed) into two parts:

- The variability explained by the model (the regression line): SSR
- The variability not explained by the model (the residuals): SSE

Sum of Squares in Total = Sum of Squares from Regression + Sum of Squared Errors

$$SST = SSR + SSE$$

$$\sum_{i=1}^n (y_i - \bar{y})^2 = \sum_{i=1}^n (\hat{y}_i - \bar{y})^2 + \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

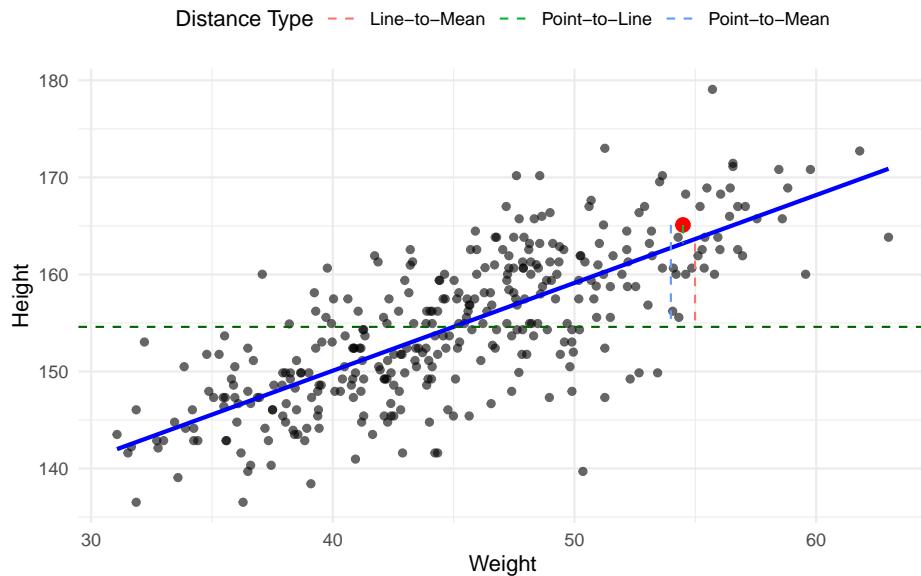
If you are interested in the details, check out this.

This video explains the concept nicely.

Let's visualize our regression result:

```
## `geom_smooth()` using formula = 'y ~ x'
```

Scatterplot with Regression Line and Differences



The blue dotted line is the distance from the mean to the point (total variance), the red dotted line is the distance from the mean to the regression line (explained variance) and the green dotted line is the distance from the regression line to the point (unexplained variance). We see that it adds up. I find this fact quite fascinating. One finds additivity by considering not deterministic values, but variances. Thank you Ronald Fisher.

See exercise 13.

3.2.5 R^2 - Coefficient of Determination

R^2 is the **amount of variance explained by the model**. You can also read Westfall 8.1.

As you can see above, the total variance (SST) of our outcome (height) can be decomposed into two parts: the variance explained by the model (SSR) and the variance not explained by the model (SSE).

Maybe the most intuitive definition of R^2 is:

$$R^2 = \frac{SSR}{SST} = \frac{SST - SSE}{SST} = 1 - \frac{SSE}{SST}$$

The value is between 0 and 1. The higher the value, the more variance is explained. But be cautious. Depending on the context, a really high R^2 is not necessarily a good thing. With the data we are working with, it could easily

hint towards an error. If we are near 1, all points in the simple linear regression model are on the line. If we are near 0, the model does not explain much of the variance and we would see “noise with no slope” in the scatterplot (exercise 4). The normal R^2 can be found in the R output under **Multiple R-squared**.

If you add a lot of variables to your regression model, you can get an R^2 of 1. The R^2 will never decrease when adding more variables. We will verify this when we have more than 2 explanatory variables. As a non-formal explanation for this: In the Sum of Squares Errors (SSE), if you add more covariates (β_2, β_3), you have more freedom to choose values that minimize the number that will be squared. Simple regression is just a special case of multiple (more than one predictor) regression with $\beta_2 = \beta_3 = \dots = 0$. Hence, you will definitely not be worse off with regards to SSE when using more covariates. A smaller SSE implies a larger SSR (sum constraint; SST remains constant) and hence a larger R^2 . If you have as many explanatory variables as data points, you can get an R^2 of 1. This is overfitting at its “best” (which we want to avoid). You would just get a value for each data point by setting all other β_i to zero and $\beta_i = \frac{y_i}{x_i}$. Since we want to find the underlying process, we want to avoid this.

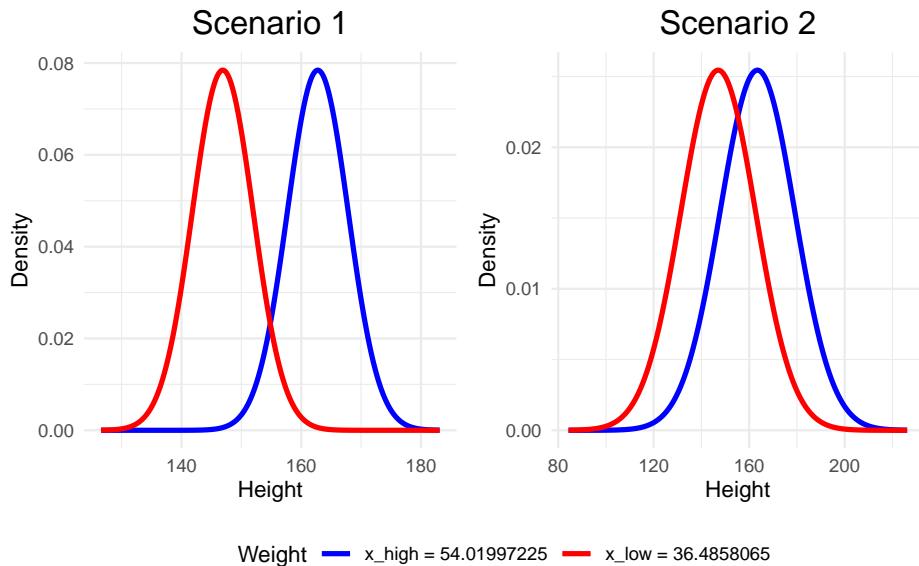
Although not perfect, one way to mitigate the influence of “too many” variables on R^2 is to use the adjusted R^2 , which can also be found in the R output (**Adjusted R-squared**).

3.2.5.1 Separating property of regression due to R^2 :

Peter Westfall explains (in Figure 8.1 of the book) how R^2 influences the separation of distributions in our simple regression model.

In our regression of *height on weight* (the order is correct, that’s how you say it), the R^2 is 0.5696. The following plot shows how “well” (i.e. precise) one can predict height if we use the 10% and 90% quantile of the weights (x_{low} and x_{high}). In both, you see the conditional distribution of height given the weight $X = x_{low}$ or $X = x_{high}$. Scenario 1 is the original model, scenario 2 is the same data with added noise (in Y-direction), which reduces R^2 to 0.13, much lower. In the right plot, the distributions have a large overlap and it is hard to distinguish between weights when seeing the height. With a very low R^2 , the height prediction does not really change and we could just as well use the mean model.

Normal Distributions for Two Scenarios



In the left plot, given $X = x_{low}$ gives a rather strongly shifted normal distribution of potentially observable heights for this weight compared to $X = x_{high}$. We would have a lower missclassification error when trying to distinguish heights of very light and very heavy people in the sample just by seeing the height. You can think of an even more extreme separation of these distributions, which would happen, when R^2 is very high or the true slope is much higher.

See also exercise 5.

An interesting way to look at R^2 is the following: Given, that one person is in the 90% quantile of the weight, the other is in the 10% quantile. What is the probability that the height of the person in the 90% quantile is higher than the height of the other person? We could calculate this relatively easy using theorems about additivity of normal distributions. Since we are all about application, we simulate this:

```
library(rethinking)
data(Howell1)
d <- Howell1
d2 <- d[d$age >= 18, ]
# Simulate heights for the two quantiles
n_sims <- 10000
set.seed(123)
mod <- lm(height ~ weight, data = d2)
summary(mod)$r.squared
```

```
## [1] 0.5696444
```

```

x_low_high <- quantile(d2$weight, probs = c(0.1, 0.9))
x_low_high

##      10%      90%
## 36.48581 54.01997

mean1 <- 113.8793936 + 0.9050291 * x_low_high[1]
mean2 <- 113.8793936 + 0.9050291 * x_low_high[2]
sd <- 5.086 # Standard deviation for both distributions
simulated_heights <- tibble(
  # conditional normal distribution according to the model
  low_heights = rnorm(n_sims, mean = mean1, sd = sd),
  high_heights = rnorm(n_sims, mean = mean2, sd = sd)
)

# Calculate the probability that the height of the person in the 90% quantile is higher
# than the height of the person in the 10% quantile
simulated_heights %>%
  mutate(higher_height = high_heights > low_heights) %>%
  dplyr::summarise(mean(higher_height))

## # A tibble: 1 x 1
##   `mean(higher_height)`<dbl>
## 1 0.987

```

In other words, we can be almost sure, that a person in the 90% quantile of the weight is taller than a person in the 10% quantile of the weight given this data set. See also exercise 11.

3.2.6 Check regression assumptions

Everytime we fit a model, we should check the assumptions above. We do this for different reasons (which will become clearer over the course). The assumptions are independent of each other. They can all be true or false or some can be true and some false (Wesftall, p.21). Chapter 4 in the book is dedicated to this topic. It is important to know that the assumptions are usually not met exactly. The question is *how badly* they are violated, not *if* they are violated. Furthermore, the assumptions refer to the data generating process, not the data itself. Thus, the evaluation of the assumptions should involve subject matter knowledge.

p-values to evaluate model assumptions are not a good idea. To quote the American Statistical Association (ASA): “By itself, a *p*-value does not provide

a good measure of evidence regarding a model or hypothesis.” Deicision-tree thinking might not be the best idea for statistical modeling.

We will not use hypothesis tests for assumptions, because

- They are never met exactly. You cannot “prove” them.
- With small sample sizes, the statistical *power* is often low.
- With large sample sizes, the smallest deviation will be “significant”.

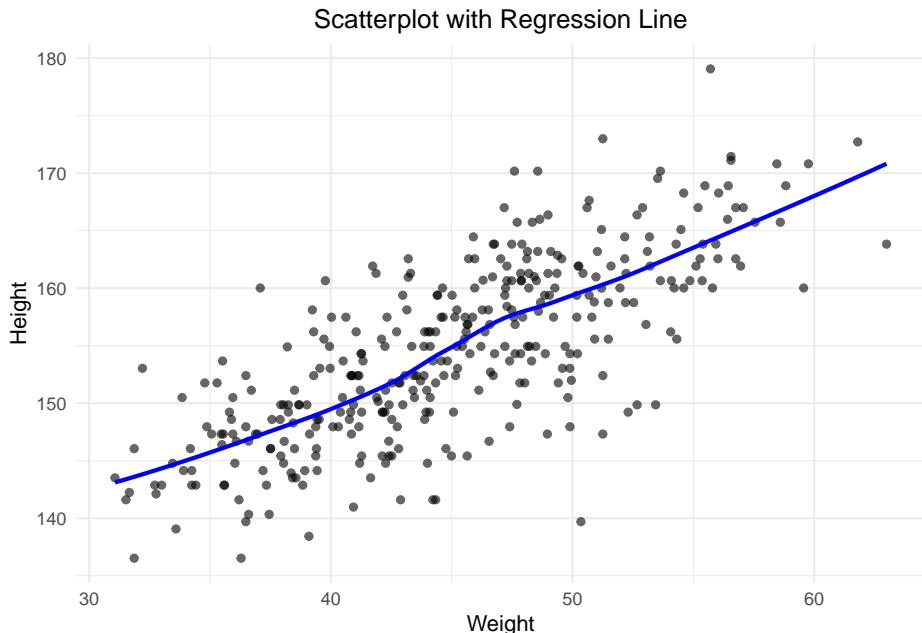
We will follow the book (chapter 4, pages 99ff) and use graphical and simulation methods to check the assumptions. As guidance, we could check them in the following order:

- Linearity
- Constant variance
- Independence
- Normality

3.2.6.1 Linearity

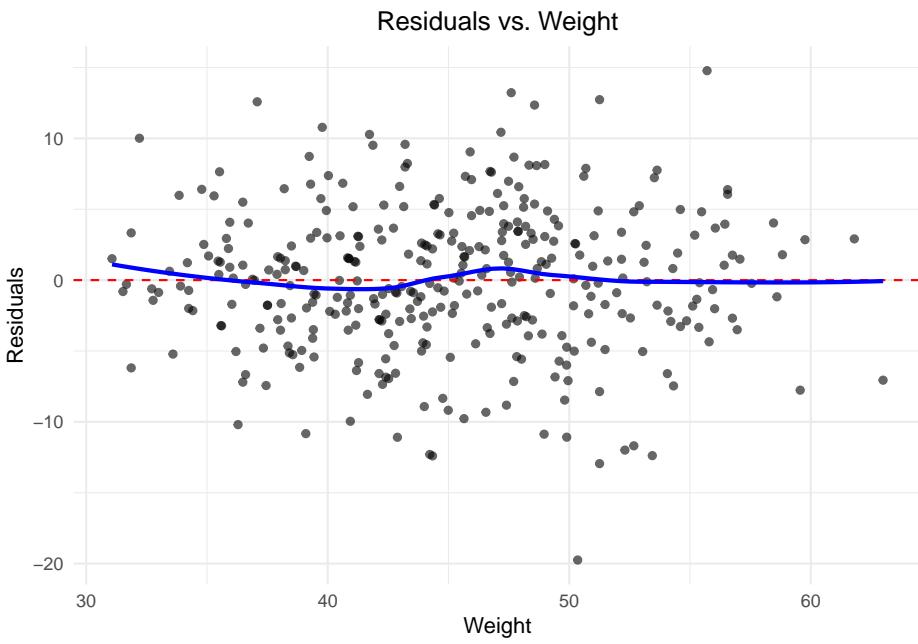
First, we **plot the data**, as we did already above. We can add a a smoothing line to the raw data as well:

```
## `geom_smooth()` using formula = 'y ~ x'
```



It looks like this relationship is describable in a linear way. No apparent curvature or patches. A refined version of the scatterplot of the raw data is The **residual scatter plot** (x_i, e_i):

```
## `geom_smooth()` using formula = 'y ~ x'
```

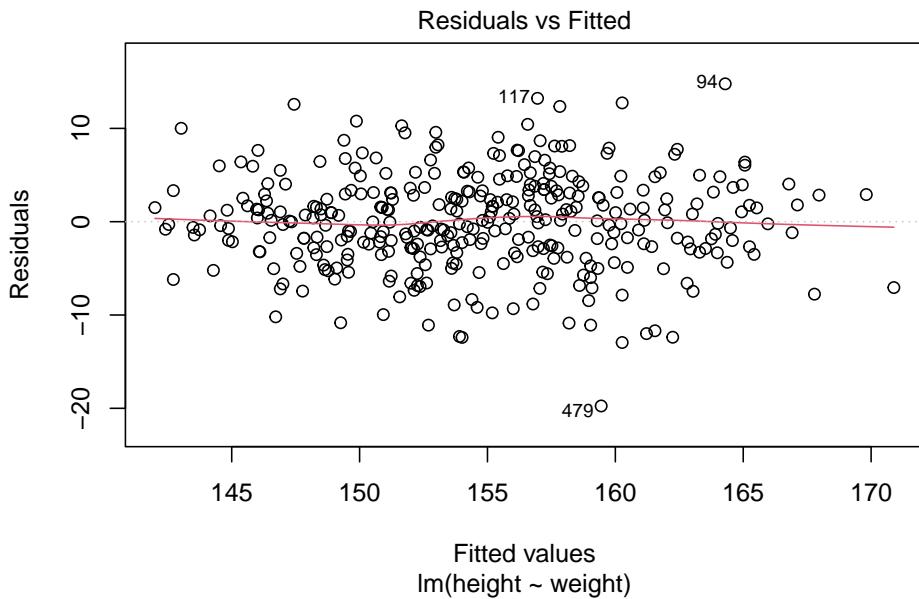


Compared to the scatterplot above, the residuals plot magnifies possible curvature. The reason is that the range of residuals is smaller than the range of the heights.

In multiple regression, we will use the (\hat{y}_i, e_i) plot, which is identical to the plot above in simple linear regression, but very helpful in multiple regression. You get the (\hat{y}_i, e_i) plot in R with `plot(mod, which = 1)`.

3.2.6.2 Constant variance

This assumption means that the variance of the residuals is constant. If it is violated, the spread around a hypothetical regression line is not constant. We look at the residual scatter plot again:



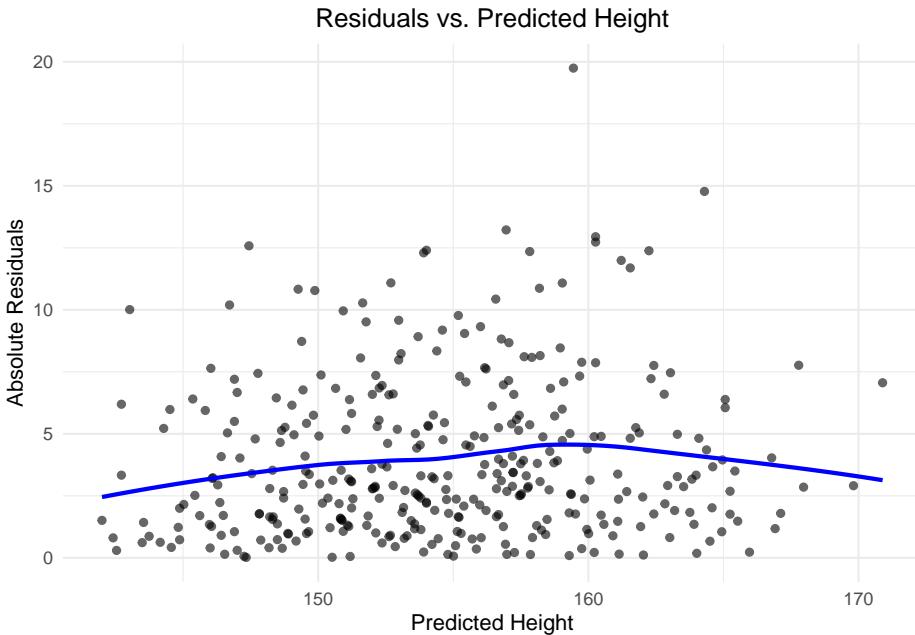
Look for changes in patterns of vertical variability. Note, that you often do not have as many data points near the end of the range of the predictor. Here are some examples of heteroskedasticity: 1, 2, 3. The above looks homoscedastic. If it was heteroskedastic, this is not a problem, we just have to model it differently (later).

As always, it is a good idea to study the variability of these plots using simulation (exercise 7).

Better for detecting heteroscedasticity is the $(\hat{y}_i, |e_i|)$ plot with a smoothing line:

```
## 'data.frame': 544 obs. of 4 variables:
## $ height: num 152 140 137 157 145 ...
## $ weight: num 47.8 36.5 31.9 53 41.3 ...
## $ age   : num 63 63 65 41 51 35 32 27 19 54 ...
## $ male  : int 1 0 0 1 0 1 0 1 0 1 ...

## `geom_smooth()` using formula = 'y ~ x'
```



This will probably not be a perfectly horizontal line, since there are less points at the end of the range of the predictor. For instance, There are less people with extreme weights in the data set.

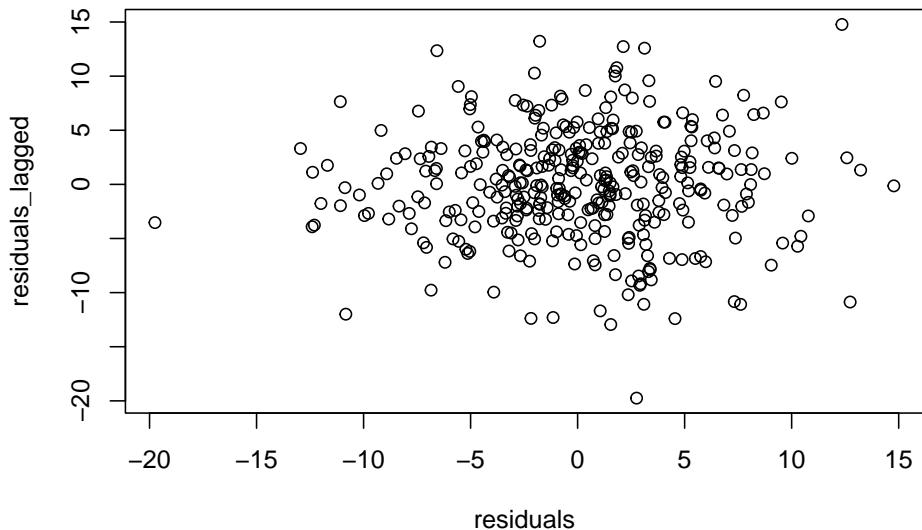
3.2.6.3 Independence, uncorrelated errors

We had an example above, where the errors were correlated. The sine curve could stem from time series data, where the x -variable is the time and the y -variable is a seasonally changing variable like temperature (purely hypothetical). In exercise 6, the values are autocorrelated: Correlated with previous values of the same variable. If we would track the body heights of persons over time, we would have an autocorrelated time series since the height of a person at time t is correlated with the height of the same person at time $t - 1$, but a little less with the height at time $t - 2$ and so on. If I am tall today, it is very likely that I was tall yesterday.

In our case of the !Kung San data, we do not have autocorrelated data.

But still, we could look at the correlations the residuals with lagged residuals. A $lag = 1$ means I compare the residuals with the ones right next to me and check if they are correlated.

```
## cor= 0.01858044
```



It least with a lag of 1, there so no large correlation between the residuals.

3.2.6.4 Normality

This assumptions states that the conditional distribution $Y|X = x$ is a normal distribution. We do not look at the normality of Y itself, since it is not a formal requirement, that Y is normally distributed. We need to assess the normality of the residuals

$$e_i = y_i - \hat{y}_i$$

I like doing this with a Q-Q plot from the R package `car` (command `qqPlot`):

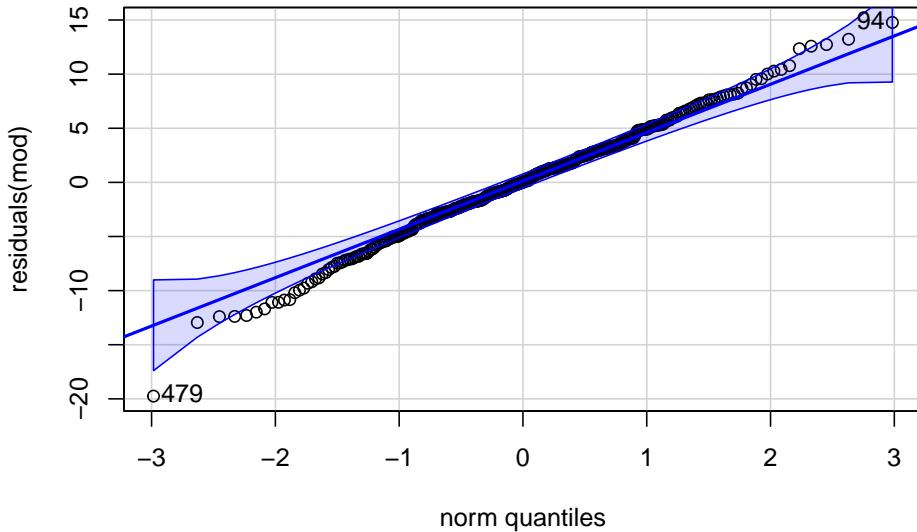
```
## Loading required package: carData

##
## Attaching package: 'car'

## The following object is masked from 'package:dplyr':
##      recode

## The following object is masked from 'package:purrr':
##      some

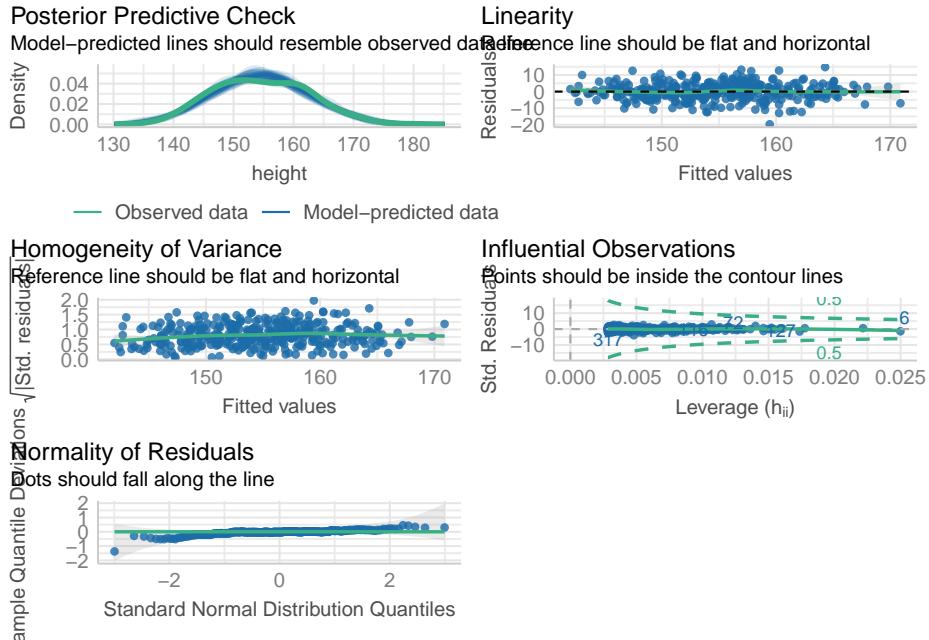
## The following object is masked from 'package:rethinking':
##      logit
```



```
## 479  94
## 317  72
```

With the 97% confidence envelope, we can see if the residuals are consistent with coming from a normal distribution. We could again use simulation to see how the Q-Q Plot changes to get a better feeling (see exercise 9).

Another convenient way to check model assumptions (for a wide class of models) is to use the `check_model` function from the `performance` package:



In this case, everything looks fine. Let's go through the plots and explain them one by one:

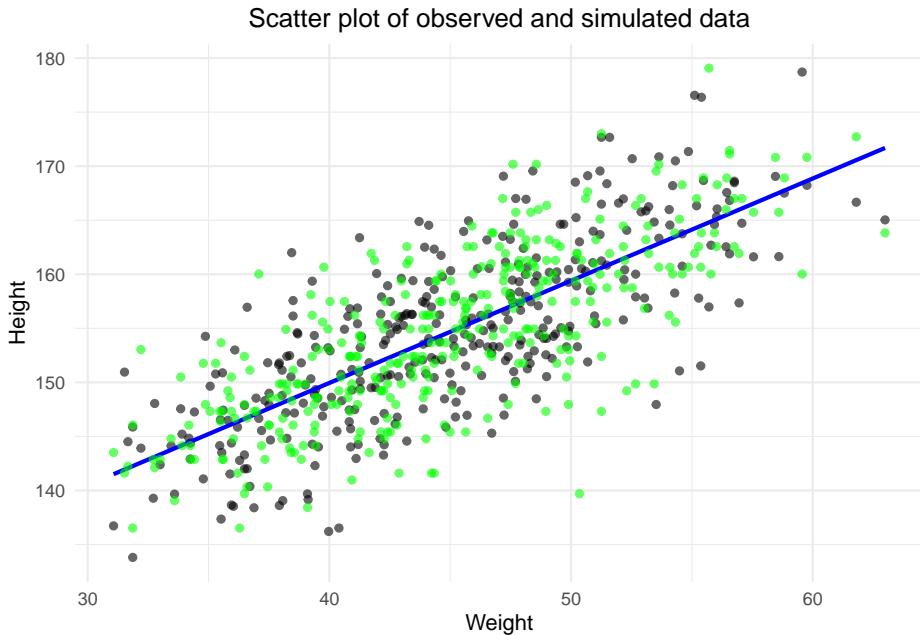
3.2.6.4.1 Posterior predictive check (upper left) Here, you create new data from the estimated model. We did that already in the Bayesian framework using `sim` from the `rstanarm` package. You could also try the command `extract.samples()` from the `rstanarm` package (Statistical Rethinking p. 196) to create new data from the Frequentist model. You can try this in our exercises sessions. Using least squares, the estimated model was:

$$\text{height}_i = 113.87939 + 0.90503 \cdot \text{weight}_i + \text{Normal}(0, 5.086)$$

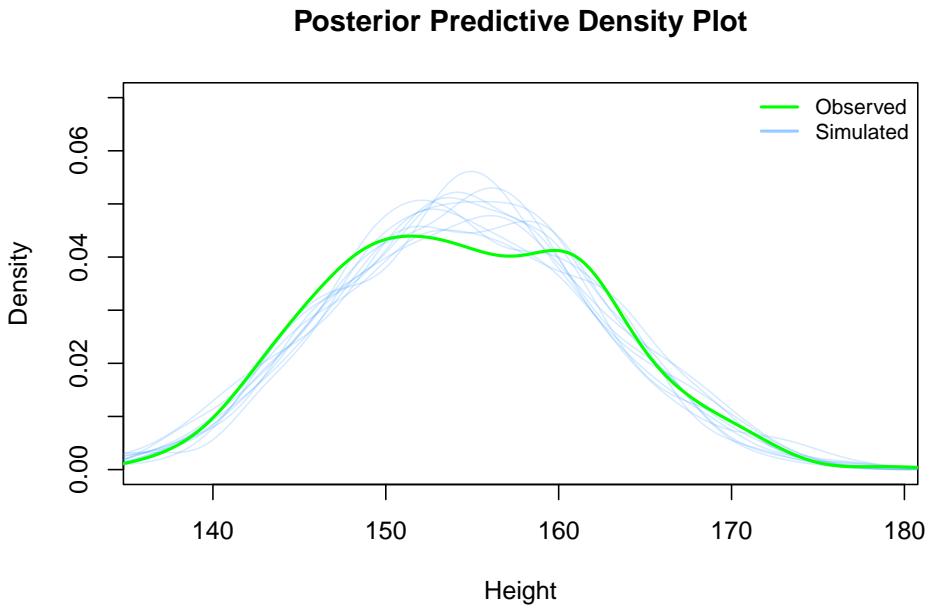
From this model, we can simulate new data (or the original weights, if we want to use fixed X) by plugging in different weights, and adding a random error from a normal distribution with mean 0 and standard deviation 5.086. Then, we can compare the distribution of the simulated data with the observed data. The blue lines in the graph are model predicted heights, the green line are the observed heights.

Let's try to replicate this plot. First, we want to simulate new data from the model. A scatter plot is also a nice way to check if the model predictions are in line with the observed data. One could repeat this process multiple times to get a feeling for the variability of the model predictions.

```
## `geom_smooth()` using formula = 'y ~ x'
```



The simulated and original data (green) fit nicely together. This lends some credibility to the model - at least in terms of prediction. And now the density plots of the model created data (blue) and the observed data (green):



As you can see, the densities of the observed and simulated data are broadly similar. One could argue that heights in the range of 150-160 cm are a bit

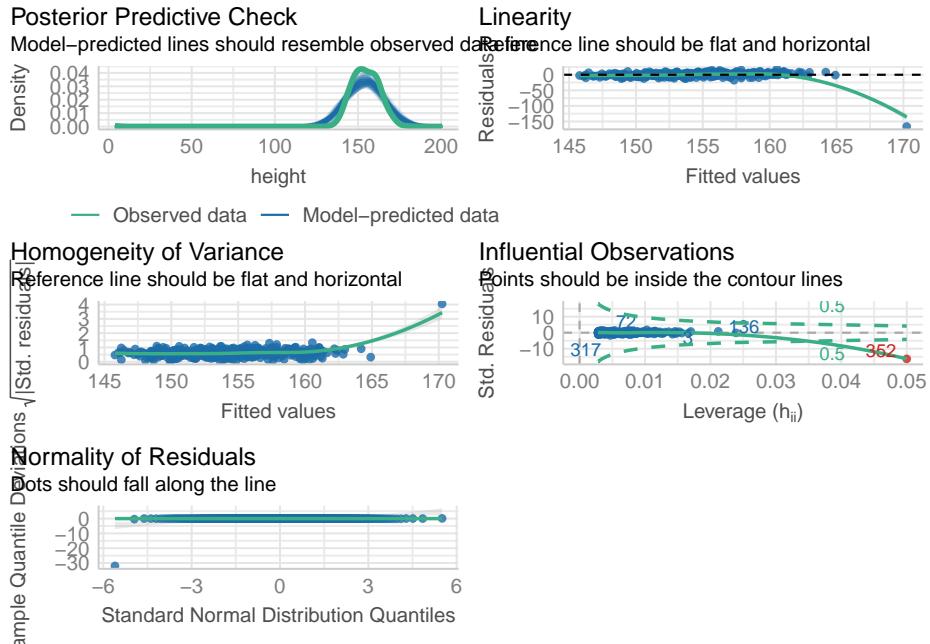
overestimated by the model. Depending on the context, this might be a problem or not. Let's accept the model for now.

3.2.6.4.2 Linear fit (upper right) This is the same plot as above: (\hat{y}_i, e_i) .

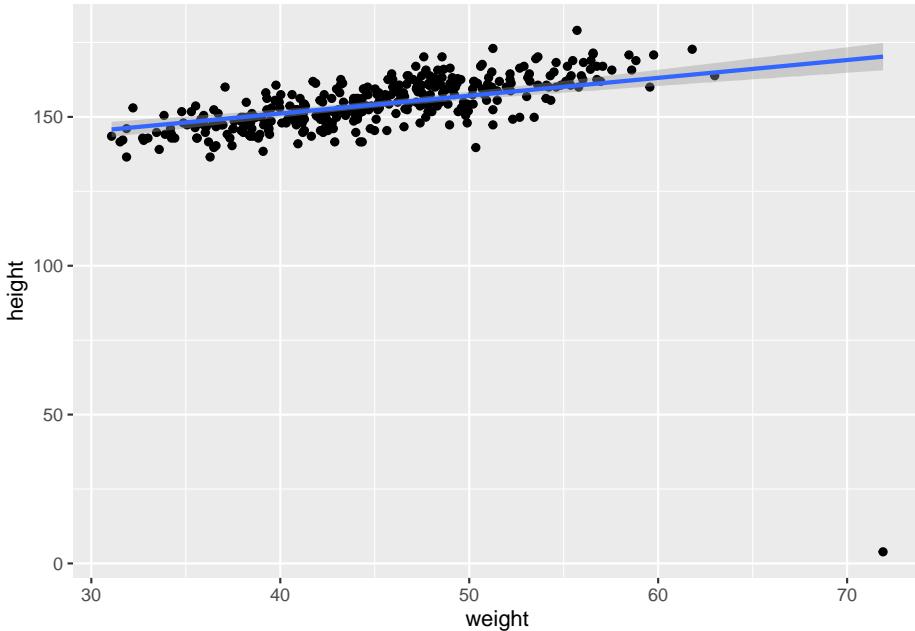
3.2.6.4.3 Homogeneity of variance (middle left) This is the same plot as above: $(\hat{y}_i, |e_i|)$.

3.2.6.4.4 Influential observations (middle right) The standard method to check for influential observations is to compute the so-called Cook's distance. This is a measure of how much leverage a single observation has on the model. We can extract the Cook's distance from the model using `cooks.distance()`. An ugly rule of thumb would be to look at observations with a Cook's distance greater than 1. In the plot, the leverage h_{ii} is on the x-axis, and the standardized residuals are on the y-axis. In short, a high leverage means that the estimated value \hat{y}_i is potentially far away from the original value y_i . The contours (dotted green lines) are the Cook's distance, in this case at a Cook's distance of 0.5. There is formula that relates the leverage (h_{ii}), the Cook's distance and the residuals. Holding Cook's distant constant at 0.5, gives you the green dotted line in the plot.

Let's create an outlier and see what happens:



```
## `geom_smooth()` using formula = 'y ~ x'
```



```
## Cook's distance = 7.025207
```

One single outlier changes the diagnostic plots. Observation 352 is clearly identified as influential. The one point changes all diagnostic plots notably. The estimates of the regression coefficients are also affected:

```
## [1] "Original Model"

## (Intercept)      weight
## 113.8793936   0.9050291

## [1] "Model with outlier"

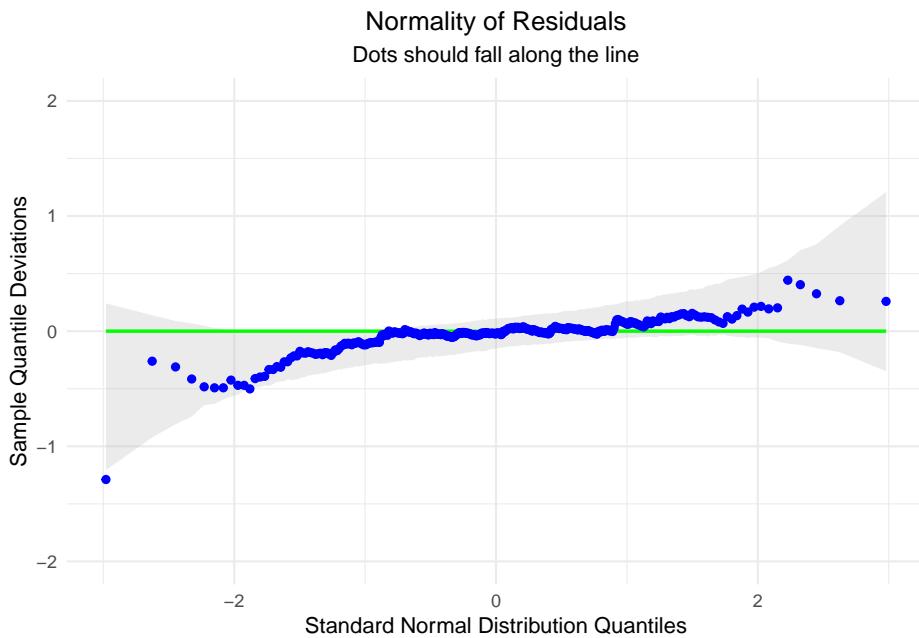
## (Intercept)      weight
## 127.172494    0.599054
```

Admitted, this is a somewhat artificial example.

3.2.6.4.5 Normality of residuals (lower left) This is basically the same plot as above, just detrended. Let's try to replicate it:

```
## 
## Attaching package: 'qqplotr'
```

```
## The following objects are masked from 'package:ggplot2':
##
##     stat_qq_line, StatQqLine
```



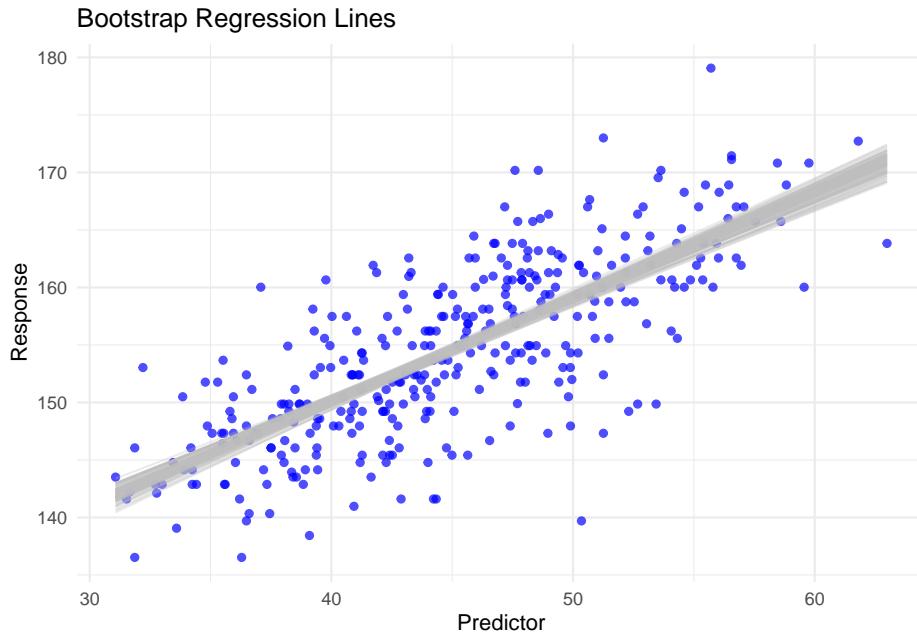
The normality assumption seems to hold. Note that in the above QQ plot, we detrended the residuals and used a bootstrap confidence band. Depending on the band type, the confidence band can be wider or narrower and include or exclude points.

After having checked the assumptions for the classical regression model and we feel comfortable with the model, we get exact confidence intervals for the effect sizes (β s) using `confint()` (p. 74 in Westfall).

Heureka! We have a model that fits the data well.

3.2.7 Bootstrap fit

In order to get a feeling for the variability of the model with regards to the predictors as well, we can bootstrap the whole data set, fit the model and draw regression lines. We create 100 bootstrap replicates of our data set `d2` by drawing with replication. For every bootstrap replicate, we fit the model and draw the regression line.



The results is very stable. Neither intercept nor slope change much.

3.2.8 Regression towards the mean

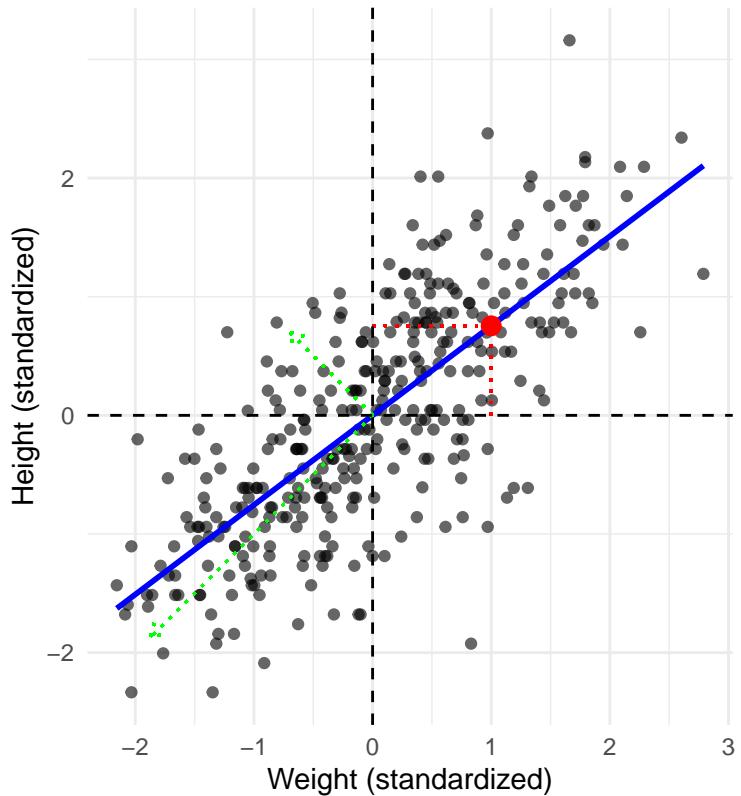
There are great explanations for “regression towards the mean” in Gelman p.58 and Westfall p.36. This video might be interesting to watch.

It describes the phenomenon that the predicted value (y) is closer to its mean than the predictor (x) to its mean. In our case this means, if the weight of a person is 1 standard deviation above the mean of body weights, the (model-)predicted height is less than 1 standard deviation above the mean of body heights, but still larger than the average height. Let's verify:

```
## predicted height= 0.7547479

## `geom_smooth()` using formula = 'y ~ x'
```

Regression towards the mean with Principal Components



As Gelman points out in his book (Figure 4.2), there is a fine detail: The regression line is not the line most people would draw through the data. They would draw a line through the main directions of variability (directions are the eigenvectors of the covariance matrix). In dotted green, you can see these directions. The regression line is a solution to another problem (as we have seen): It minimizes the sum of squared residuals, which are defined via the **vertical** distances to the line. See exercise 8.

3.2.9 Random X vs. fixed X

A detail that is not mentioned often in introductory statistics courses is the question of whether the predictor variable X is random or fixed. In our case, we did not specify the weights of the people in the !Kung San data set. *Observational* data was collected and we have no control over the weights. In an *experimental* setting, we could have controlled the weights of the people. We could have for instance only included people with weights at certain steps (50 kg, 60 kg, 70 kg). In the latter case, we could consider X as fixed. In the former case, we consider X as random. If we would draw another sample from the population,

we would get different weights again. Further reading: Westfall 1.5.

There are many more details we could look into, but we wanted to give a first practical introduction to regression analysis with less emphasis on theory and proofs behind it.

We want to make clear what the confidence intervals from `confint` of a model mean. For this see exercise 14.

3.3 Exercises

[E] Easy, [M] Medium, [H] Hard

(Some) solutions to exercises can be found in the git-repo here.

3.3.1 [E] Exercise 1

In the model from above:

$$\begin{aligned} h_i &\sim \text{Normal}(\mu_i, \sigma) \\ \mu_i &\sim \alpha + \beta(x_i - \bar{x}) \\ \alpha &\sim \text{Normal}(171.1, 20) \\ \beta &\sim \text{Normal}(0, 10) \\ \sigma &\sim \text{Uniform}(0, 50) \end{aligned}$$

- What is the expected height when $x_i = \bar{x}$?
- What is the expected height when x_i changes by 1 unit?

3.3.2 [E] Exercise 2

Look at the marginal distributions of the parameters in the Bayesian model.

- Plot the posterior distribution of all 3 parameters.
- Include in the plot a 99% credible interval (HDI).

3.3.3 [M] Exercise 3

Go to the coefficient estimates in the simple linear regression setting above (Fit the model) in the classical framework.

- Create an R file to simulate the simple linear regression model.
- Change your input parameters and see how the estimates change.
- Does this make sense with respect to the estimates given, specifically with respect to β_1 ?

3.3.4 [M] Exercise 4

Verify the statement above in the text for high and low values of R^2 .

3.3.5 [M] Exercise 5

Verify with simulation in R that the separation of the distributions in the simple linear regression model improves if the true (but usually unknown) slope increases.

3.3.6 [H] Exercise 6

Go to the model assumptions in the classical regression model (Model Assumptions). - Use the code from github to recreate the regression model with the sine-curve. - Check the independence assumption as described. Look at the residuals, when $X = 2$ and $X = 4.5$. You can get those by filtering residuals that are > 0.5 and < 0.5 .

3.3.7 [M] Exercise 7

- Simulate data from the regression of heights on weights in our !Kung San data set.
- Draw the \hat{y}_i, e_i plot.
- Draw the $\hat{y}_i, |e_i|$ plot.
- Repeat the simulation and look at the variability of the plot.

3.3.8 [E] Exercise 8

- Go to p.36 in Westfall's book and read Appendix A.
- Pay close attention to the explanation about regression toward the mean.

3.3.9 [M] Exercise 9

Go to the resuials above where we tested the normality assumption.

- Calculate mean and standard deviation from the residuals of the model that regresses height on weight.
- Simulate from a normal distribution using these parameters.
- Get a feeling how the QQ plot changes by drawing the QQ plot repeatedly.

3.3.10 [M] Exercise 10

Using our !Kung San data,

- show that the regression of height on weight (`lm(height ~ weight)`) is not the same as the regression of weight on height (`lm(weight ~ height)`).
- Draw both regression lines in one diagram.
- Can you simulate data where the two regressions deliver (almost) identical results?
- Explain why the results differ and what consequences this would have for a research question. Which question do I answer with each?

3.3.11 [E] Exercise 11

Go back to the section about R^2 and the separation of the distributions. How would the probability that a person in the 90% quantile of the weight is taller than a person in the 10% quantile of the weight change if you change

- the true slope of the regression line
- the true σ , i.e. if you add more noise and have a lower R^2 ?

3.3.12 [M] Exercise 12

Go back to the Bayesian simple regression model for height above.

- Standardize weights and heights in the data set `d2`.
- Estimate the regression model as we did in the section using `quap`.
- Plot the posterior distribution using `plot(model)`.
- Interpret the regression coefficient β .

3.3.13 [M] Exercise 13

Go back to the ANOVA section.

- Calculate SSE, SST and SSR for the regression of height on weight.
- How many degrees of freedom does each term have?

3.3.14 [M] Exercise 14

When we fit a simple linear regression model (Frequentist), we get confidence intervals for the coefficients using the command `confint`. As stated previously, these are frequentist CIs. That means, if we draw data from the true (but usually unknown) data generating process, we would expect that the true value of the coefficient (intercept and slope) lies in the CI in X% of the cases.

- Simulate data from the true data generating process of the simple linear regression model.
- Fit the model and get the confidence intervals for the coefficients.
- Repeat this process 1000 times and see how many times the true value of the coefficients lies in the CI.

Changing the distribution of the covariate x should not change the results, since we only get a different amount of data points across the range of x .

- Verify this statement by changing the distribution of x .

Chapter 4

Multiple Linear Regression

So far, we have dealt with the simple mean model and the model with one predictor in the Bayesian and Frequentist framework. We will now add another predictor and subsequently an interaction term to the model. Finally, we will add more than two predictors to the model.

If you feel confused at any point: As Richard McElreath repeatedly says: This is normal, it means you are paying attention. I also refer to the great Richard Feynman.

4.1 Linear Regression with 2 predictors in the Bayesian Framework

4.1.1 Meaning of “linear”

What is a linear model? The term “linear” refers to the relationship of the predictors with the dependent variable (or outcome). The following model is also linear:

$$height_i = \beta_0 + \beta_1 x_i + \beta_2 x_i^2$$

The model is linear in the parameters $\beta_0, \beta_1, \beta_2$ but not in the predictors x_i . The term x_i^2 is ok, since the heights are just sums of multiples of the predictors (which can be nonlinear). This model is not a linear model anymore:

$$height_i = \beta_0 + \beta_1 x_i + e^{\beta_2 x_i^2}$$

β_2 is now the exponent of e . It would also not be linear, if the coefficients are in a square root or in the denominator of a fraction, or in a sine or in a logarithm. You get the idea.

Here is an easy way to check if the model is linear: If I change the predictor value (i.e., the value of x_i , x_i^2 or whatever your predictor is) by one unit, the change in the (expected value of the) dependent variable is the coefficient in front of the predictor (β_i).

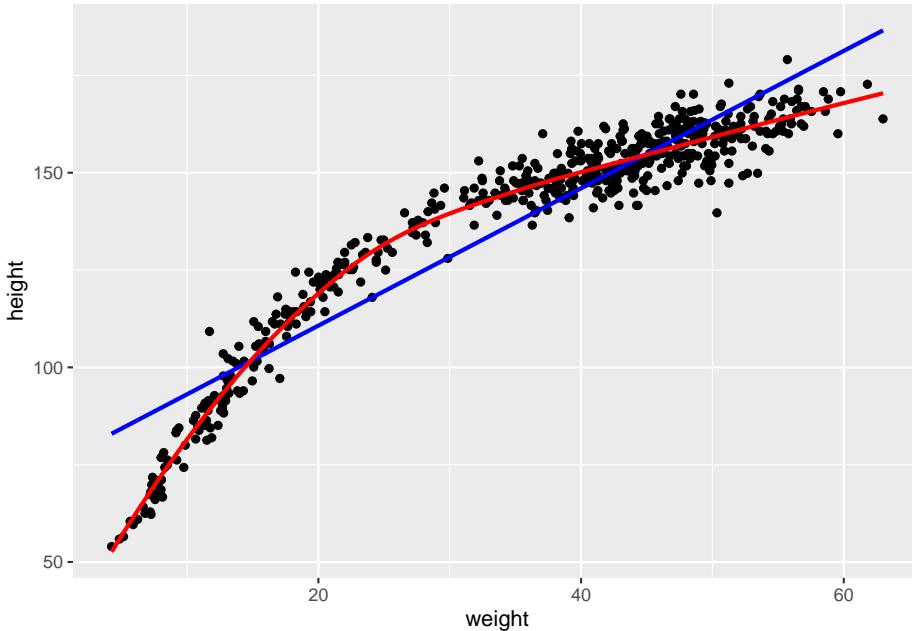
4.1.2 Adding a transformed predictor to the model

Around 4.5. in the book Statistical Rethinking there is a lineare regression using a quadratic term for weight. It is a principle, called the “**variable inclusion principle**”, that we always include the lower order terms when fitting a model with higher order terms. See Westfall, p. 213. If we do not include the lower order terms, the coefficient does not measure what we want it to measure (curvature in our case). For instance, if we want to model a quadratic relationship (parabola) between weight and height, we also have to include the linear term for weight (x_i). Since we do not assume the relationship between weight and height to be linear but quadratic (which is a polynomial of degree 2), we call this a polynomial regression. This video could be instructive. One has to be careful with fitting polynomials to data points since the regression coefficients can become quite large. Using a polynomial of high degree implies to have a lot of parameters to estimate. Increasing the degree of the polynomial increases R^2 but also the risk of overfitting. (see Statistical Rethinking p. 200). So this is - of course - not the final solution to regression problems.

This time, lets look at the whole age range of data from the !Kung San people.

```
library(rethinking)
library(tidyverse)
data(Howell1)
d <- Howell1
d %>% ggplot(aes(x = weight, y = height)) +
  geom_point() +
  geom_smooth(method = "lm", se = FALSE, color = "blue") +
  geom_smooth(method = "loess", se = FALSE, color = "red")

## `geom_smooth()` using formula = 'y ~ x'
## `geom_smooth()` using formula = 'y ~ x'
```



It would not be a good idea to fit a linear trend through this data, because we would not capture the relationship adequately. The red line is a loess smoothing line which is often used to capture non-linear relationships. The blue line is the usual line from classic linear regression (from the previous chapter). Which one describes the data more accurately? In this case it is obvious, a non-linear relationship is present and it might be a good idea to model it. Modeling the relationship with a linear trend leads to bad residuals with structure. We will demonstrate this in the frequentist setting. Unfortunately, in more complex settings, with more predictors, it is not always so easy to see.

This time, we use the mean for the prior from the book (178cm). The model equations are (see exercise 2):

$$\begin{aligned}
 h_i &\sim \text{Normal}(\mu_i, \sigma) \\
 \mu_i &= \alpha + \beta_1 x_i + \beta_2 x_i^2 \\
 \alpha &\sim \text{Normal}(178, 20) \\
 \beta_1 &\sim \text{Log-Normal}(0, 1) \\
 \beta_2 &\sim \text{Normal}(0, 1) \\
 \sigma &\sim \text{Uniform}(0, 50)
 \end{aligned}$$

The prior for β_1 is log-normal, because we can reasonably assume the overall linear trend is positive. The prior for β_2 is normal, because we are not so sure. If we thought back to our school days to the topic of “curve discussion” or

parabolas, we could probably also assume that β_2 is negative. But, data will show.

How can we interpret the model equations? The model assumes that the **expected** height μ_i of a person i depends non-linearly (quadratically) on the weight x_i of the person. We are in the business of mean-modeling. The prior for σ is uniform as before. The prior for α is normal with mean 178 and standard deviation 20 because this is what we can expect from body heights in our experience.

Let's **fit the model**:

We standardize the weight again and add the squared weights to the data set. Standardizing the predictors is a good idea, especially in polynomial regression since squares and cubes of large numbers can get huge and cause numerical problems.

Let's fit the model with the quadratic term for weight:

```
# Standardize weight
d$weight_s <- (d$weight - mean(d$weight)) / sd(d$weight)
# Square of standardized weight
d$weight_s2 <- d$weight_s^2
m4.1 <- quap(
  alist(
    height ~ dnorm(mu, sigma),
    mu <- a + b1*weight_s + b2*weight_s^2,
    a ~ dnorm(178, 20),
    b1 ~ dnorm(0, 10),
    b2 ~ dnorm(0, 10),
    sigma ~ dunif(0, 50)
  ), data = d)
precis(m4.1)

##           mean        sd      5.5%     94.5%
## a     146.672739 0.3736465 146.075580 147.269898
## b1    21.397637 0.2898827  20.934348 21.860925
## b2   -8.419933 0.2813308 -8.869554 -7.970312
## sigma 5.750550 0.1743749   5.471865  6.029235
```

β_2 is indeed negative. We get our **joint distribution** of the **four model parameters**. Let's look at the fit using the mean estimates of the posterior distribution:

```
# Summarize the model parameters
model_summary <- precis(m4.1)
params <- as.data.frame(model_summary)
```

```

# Extract parameter values
a <- params["a", "mean"]           # Intercept
b1 <- params["b1", "mean"]         # Coefficient for standardized weight
b2 <- params["b2", "mean"]         # Coefficient for squared standardized weight

# Generate a sequence of standardized weights for the fitted curve
weight_seq <- seq(min(d$weight_s), max(d$weight_s), length.out = 200)

# Calculate the fitted values using the quadratic equation
height_fitted <- a + b1 * weight_seq + b2 * weight_seq^2

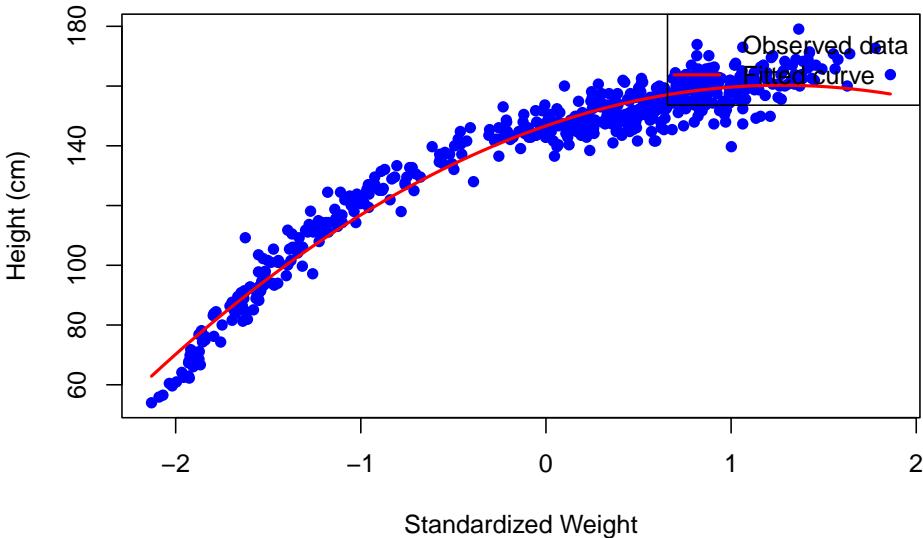
# Plot the scatterplot
plot(d$weight_s, d$height, pch = 16, col = "blue",
      xlab = "Standardized Weight", ylab = "Height (cm)",
      main = "Scatterplot with Fitted Curve (Standardized Weight)")

# Add the fitted curve
lines(weight_seq, height_fitted, col = "red", lwd = 2)

# Add a legend
legend("topright", legend = c("Observed data", "Fitted curve"),
       col = c("blue", "red"), pch = c(16, NA), lty = c(NA, 1), lwd = 2)

```

Scatterplot with Fitted Curve (Standardized Weight)



This fits much better than the linear model without the quadratic term. In the book, there is also a polynomial regression with a cubic term for weight. Maybe

this fits even better (see exercise 1).

4.1.3 Adding another predictor to the model

Since the !Kung San data set has already such a high R^2 with the quadratic term (and possibly higher with the cubic term), we will use the created data set from below in the frequentist setting to estimate the coefficients of the model with two predictors.

We use rather uninformative priors and fit the model using `quap`:

```
library(rethinking)
set.seed(123)
n <- 100
X1 <- rnorm(n, 0, 5)
X2 <- rnorm(n, 0, 5)
Y <- 10 + 0.5 * X1 + 1 * X2 + rnorm(n, 0, 2) # true model
df <- data.frame(X1 = X1, X2 = X2, Y = Y)

# fit model
m4.2 <- quap(
  alist(
    Y ~ dnorm(mu, sigma),
    mu <- a + b1*X1 + b2*X2,
    a ~ dnorm(10, 10),
    b1 ~ dnorm(0, 10),
    b2 ~ dnorm(0, 10),
    sigma ~ dunif(0, 50)
  ), data = df)
precis(m4.2)

##           mean        sd      5.5%     94.5%
## a     10.2700227 0.18934442 9.9674137 10.5726316
## b1    0.4467278 0.04131434 0.3806995  0.5127561
## b2    1.0095082 0.03899992 0.9471788  1.0718377
## sigma 1.8738833 0.13250803 1.6621099  2.0856567
```

4.1.3.1 Checking model assumptions

Andrew Gelman mentions in some of his talks (see here for more details) that many Bayesians he met do not check their models, since they reflect subjective probability. As I said in the introduction, one should not be afraid to check model predictions against the observed and probably new data. If a model for predicting BMI performs much worse on a new data set, we can probably

conclude that the model does not reflect the general relationship between the predictors and the dependent variable. We **do not ask** the question if a model is true or false, but if it is useful or how badly the model assumptions are violated.

For further, more detailed information on model checking, refer to chapter 6 of Gelman's book.

Anyhow, we plot two posterior predictive checks here. We test the model within the same data set. In order to do this, we create new observations by drawing from the posterior distribution and compare these with the acutally observed values. This is called **posterior predictive checks**.

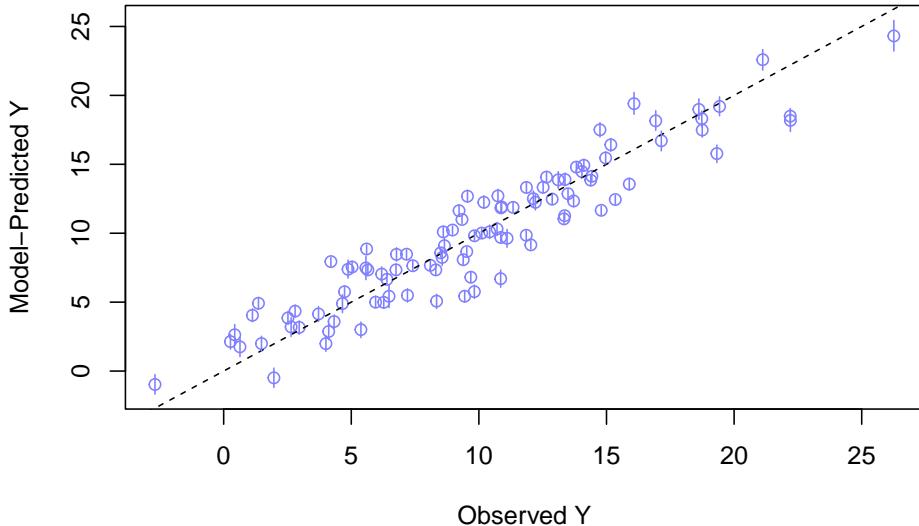
First, we plot the observed Y values against the predicted Y values ($= \hat{Y}$) from the model (as in Statistical rethinking, Chapter 5). Although these practically never lie on the line $y = x$, they should be sufficiently close to it. We could also compare these two plots with the mean-model (see exercise 6).

```
# 1) Posterior predictive checks Y vs Y_hat
# see Statistical Rethinking p 138.
# call link without specifying new data
# so it uses the original data
mu <- link(m4.2)

# summarize samples accross cases
mu_mean <- apply(mu, 2, mean)
mu_PI <- apply(mu, 2, PI, prob = 0.89)

# simulate observations
# again, no new data, so uses original data
D_sim <- sim(m4.2, n = 1e4)
D_PI <- apply(D_sim, 2, PI, prob = 0.89)

plot(mu_mean ~ df$Y, col = rangi2, ylim = range(mu_PI),
      xlab = "Observed Y", ylab = "Model-Predicted Y")
abline(a = 0, b = 1, lty = 2)
for(i in 1:nrow(df)) lines(rep(df$Y[i], 2), mu_PI[,i], col = rangi2)
```



As we can see, the model fits the data quite well. The points are close to the dashed line ($y = x$). No under- or overestimation is visible. The model seems to capture the relationship between the predictors X_1 and X_2 and the dependent variable Y quite well - at least in a predictive sense. If there were patches of data points above or below the dashed line, we would probably have to reconsider the model definition and think about why these points are not captured by the model.

Next, we plot the posterior predictive plots analog to the upper left in the `check_model` output.

```
library(scales) # For the alpha function to adjust transparency

##
## Attaching package: 'scales'

## The following object is masked from 'package:purrr':
##      discard

## The following object is masked from 'package:readr':
##      col_factor

# 2) Posterior predictive densities
# Simulate observations using the posterior predictive distribution
D_sim <- sim(m4.2, n = 1e4) # Generate 10,000 simulated datasets
```

```

# Calculate densities for all samples
densities <- apply(D_sim, 1, density)

# Find the maximum density value for setting the y-axis limits
max_density <- max(sapply(densities, function(d) max(d$y)))

# Create the density plot with predefined ylim
plot(NULL, xlim = range(df$Y), ylim = c(0, max_density),
      xlab = "Y", ylab = "Density",
      main = "Comparison of Observed and Predicted Densities")

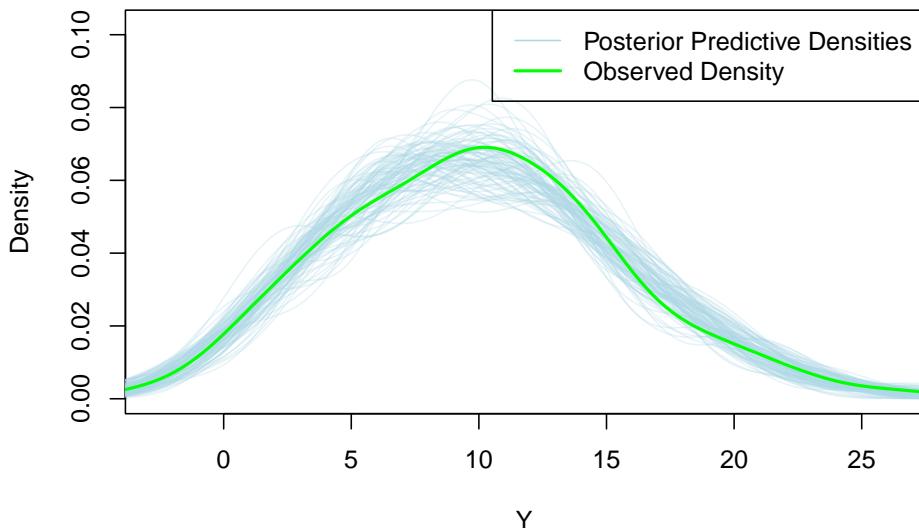
# Add 100 posterior predictive density lines
set.seed(42) # For reproducibility
n_lines <- 100
samples <- sample(1:1e4, n_lines) # Randomly sample 100 posterior predictive datasets
for (s in samples) {
  lines(density(D_sim[s, ]), col = alpha("lightblue", 0.3), lwd = 1)
}

# Add the density line for the observed Y values
obs_density <- density(df$Y)
lines(obs_density$x, obs_density$y, col = "green", lwd = 2)

# Add legend
legend("topright", legend = c("Posterior Predictive Densities", "Observed Density"),
       col = c("lightblue", "green"), lty = 1, lwd = c(1, 2))

```

Comparison of Observed and Predicted Densities



The light blue lines show distributions of model predicted Y values. The green line shows the distribution of the observed Y values. As we can see, there seem to be no systematic differences between the observed and predicted values. The model seems to capture the relationship well. If we see systematic deviations here, we need to reconsider the model definition.

Example: If you want to predict pain (Y variable) and you have a lot of zeros (pain-free participants) you will probably see a discrepancy between the observed and predicted values in this plot. What could you do? You could use a two step process (model the probability that a person is pain-free and then model the pain intensity for the people who have pain) or use a different model (like a zero-inflated model).

Note that we did not explicitly assume normally distributed errors in the model definition above, so we won't check this here but in the Frequentist framework below.

4.2 Linear regression with 2 predictors in the Frequentist Framework

4.2.1 Adding a transformed predictor to the model

No, let's fit the same model as above in the Frequentist framework.

The model is:

$$height_i = \alpha + \beta_1 weight_i + \beta_2 weight_i^2 + \varepsilon_i$$

whereas

$$\varepsilon_i \sim N(0, \sigma)$$

And if you build the expectation on both sides for fixed $weight_i$, you get:

$$\mathbb{E}(height_i | weight_i) = \alpha + \beta_1 weight_i + \beta_2 weight_i^2$$

The last line means, the expected height of a person given a certain weight depends quadratically on the weight. The error term ε_i is on average zero, hence it goes away here. Remember the law of large numbers: The sample mean $\bar{\varepsilon}_i$ approaches the expected value $\mathbb{E}(\varepsilon_i) = 0$ as the sample size increases. If you drew many samples and average over the error terms, the average will approach zero. The weights are considered fixed and therefore do not change when building the expectation.

We are looking for fixed, but unknown, parameters α, β_1, β_2 and σ . The fixed σ indicates that the observations *wiggle* around the expected value equally strong not matter which weight we have. This is called **homoscedasticity**.

Let's fit the model using the `lm` function in R which uses least squares to estimate the parameters. At this point I could torture you with matrix algebra and show you the normal equations for linear regression, but I will spare you for now. Note that the least squares algorithm for fitting the curve works for all kinds of functional forms. For example, we could also fit an exponential curve using the same technique (see exercise 9).

```
# scale weight
d$weight_s <- scale(d$weight)
# Fit the model
m4.2 <- lm(height ~ weight_s + I(weight_s^2), data = d)
summary(m4.2)

##
## Call:
## lm(formula = height ~ weight_s + I(weight_s^2), data = d)
##
## Residuals:
##      Min       1Q   Median       3Q      Max 
## -19.9689  -3.9794   0.2364   3.9262  19.5182 
## 
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)    
## (Intercept) 146.6604    0.3748 391.30  <2e-16 ***
```

```

## weight_s      21.4149     0.2908   73.64 <2e-16 ***
## I(weight_s^2) -8.4123     0.2822  -29.80 <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 5.766 on 541 degrees of freedom
## Multiple R-squared:  0.9565, Adjusted R-squared:  0.9564
## F-statistic:  5952 on 2 and 541 DF,  p-value: < 2.2e-16

mean(d$height)

## [1] 138.2636

confint(m4.2, level = 0.94)

##                   3 %         97 %
## (Intercept) 145.954018 147.366836
## weight_s     20.866788  21.962979
## I(weight_s^2) -8.944251 -7.880337

```

See `?I` in R. This command is used so that R knows that it should treat the “`^2`” as “square” and not as formula syntax. We could also create a new variable as before. Whatever you prefer.

4.2.1.1 Interpretation of output and coefficients

- The intercept α is the **model-predicted height** of a person of **average weight** ($weight_s = 0$ for a person of average weight). Note that this is not equal to the average height (138.2636cm) of the people in the data set.
- The residuals have range from -19.97 to 19.51 . So, the model maximally overestimates the heights by 19.97 cm and underestimates by 19.51 cm. These numbers are plausible when you look at the scatterplot with the fitted curve.
- The coefficients β_1 and β_2 agree with the Bayes estimates. Specifically, β_2 is non-zero indicating curvature. You cannot directly interpret the coefficients as in the non-quadratic case since, for instance, you cannot change $weight^2$ by one unit and hold $weight$ constant at the same time. Refer to Peter Westfall’s book section 9.1. for all the details.
- If you like p -values: All the hypotheses that the coefficients are zero are rejected. The p -values are very small. The values of the test statistics can not be explained by chance alone. On the other hand, for at least β_1 and the global test this is not a surprise when you look at the

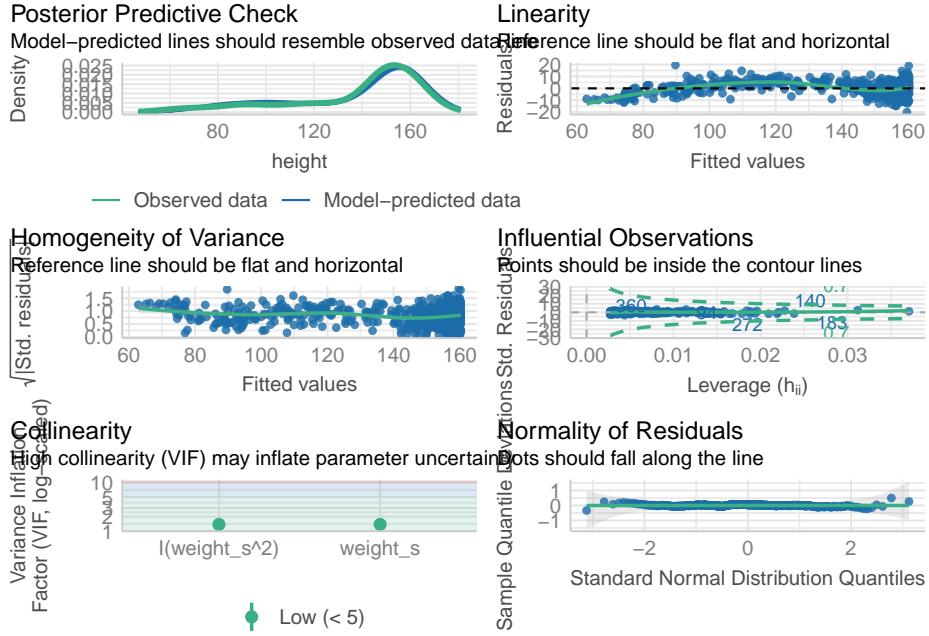
scatterplot. After having fit many models, you would have guessed that all three parameters are solidly non-zero. The intercept is not zero since a person of average weight probably has non-zero height. β_1 is non-zero since you can easily imagine a linear trend line with positive slope going through the data, and β_2 is non-zero since there is clearly (non-trivial) curvature in the scatterplot.

- The R^2 is a whopping 0.96 which could be a sign of overfitting, but in this case we conclude that the true relationship is captured rather well. Overfitting would occur if our curve would wiggle around the data points, so we would fit the data too much to the noise in the data than the underlying trend.

4.2.1.2 Checking model assumptions

```
check_model(m4.2)
```

```
## Some of the variables were in matrix-format - probably you used
## `scale()` on your data?
## If so, and you get an error, please try `datawizard::standardize()` to
## standardize your data.
## Some of the variables were in matrix-format - probably you used
## `scale()` on your data?
## If so, and you get an error, please try `datawizard::standardize()` to
## standardize your data.
```



If we want to be perfectionists, we could remark that (upper right plot) in the lower fitted values the residuals are more negative, meaning that the model overestimates the heights in this region. In the middle region the model underestimates a bit and we can see a positive tendency in the residuals. Apart from that, the diagnostic plots look excellent.

4.2.2 Adding another predictor to the model

Now, we add another predictor to the model. We use X_1 and X_2 **simultaneously** to predict Y . We are now in the lucky situation that we can still visualize the situation in 3D. The regression line from simple linear regression becomes a plane. The vertical distances between the data points and the plane are the residuals. See here or here at the end for examples. Minimizing the sum of the squared errors gives again the estimates for the coefficients.

For demonstration purposes, we can **create data ourselves** with known coefficients. This is the same as above. This is the true model, which we usually do not know:

$$Y_i = \beta_0 + \beta_1 X_{1i} + \beta_2 X_{2i} + \varepsilon_i$$

$$\varepsilon_i \sim N(0, \sigma^2)$$

$$\mathbb{E}(Y_i | X_1 = x_1; X_2 = x_2) = \beta_0 + \beta_1 x_1 + \beta_2 x_2$$

$$i = 1 \dots n$$

for example:

$$\begin{aligned} Y_i &= 10 + 0.5 \cdot X_{1i} + 1 \cdot X_{2i} + \varepsilon_i \\ \varepsilon_i &\sim N(0, 5) \\ \mathbb{E}(Y_i | X_1 = x_1; X_2 = x_2) &= 10 + 0.5x_1 + 1x_2 \\ i &= 1 \dots n \end{aligned}$$

According to the model, the conditional expected value of Y_i given $X_1 = x_1$ and $X_2 = x_2$ is a linear function of x_1 and x_2 . Note, that small letters are realized values of random variables. Also note, that in the expectation the error term goes away, since $\mathbb{E}(\varepsilon_i) = 0$.

- If X_1 increases by one unit, Y increases by 0.5 units on average (in expectation).
- If X_2 increases by one unit, Y increases by 1 unit on average (in expectation).
- If X_1 and X_2 are zero, Y is 10 on average (in expectation).

Why in expectation? Because there is still the error term which makes the whole thing random! We can see that an increase in X_1 does not influence the relationship between X_2 and Y . Hence, there is **no interaction** between X_1 and X_2 with respect to Y .

Now let's draw 100 points from this model, fit the model and add the plane:

```
library(plotly)

set.seed(123)
n <- 100
X1 <- rnorm(n, 0, 5)
X2 <- rnorm(n, 0, 5)
Y <- 10 + 0.5 * X1 + 1 * X2 + rnorm(n, 0, 2)
d <- data.frame(X1 = X1, X2 = X2, Y = Y)

# Fit the model
m4.3 <- lm(Y ~ X1 + X2, data = d)
summary(m4.3)

## 
## Call:
## lm(formula = Y ~ X1 + X2, data = d)
## 
## Residuals:
```

```

##      Min      1Q Median      3Q      Max
## -3.7460 -1.3215 -0.2489  1.2427  4.1597
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)
## (Intercept) 10.27013   0.19228  53.41 <2e-16 ***
## X1          0.44673   0.04195  10.65 <2e-16 ***
## X2          1.00952   0.03960  25.49 <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1.903 on 97 degrees of freedom
## Multiple R-squared:  0.8839, Adjusted R-squared:  0.8815
## F-statistic: 369.1 on 2 and 97 DF,  p-value: < 2.2e-16

# Create a grid for the plane
X1_grid <- seq(min(d$X1), max(d$X1), length.out = 20)
X2_grid <- seq(min(d$X2), max(d$X2), length.out = 20)
grid <- expand.grid(X1 = X1_grid, X2 = X2_grid)

# Predict the values for the grid
grid$Y <- predict(m4.3, newdata = grid)

# Convert the grid into a matrix for the plane
plane_matrix <- matrix(grid$Y, nrow = length(X1_grid), ncol = length(X2_grid))

# Create the interactive 3D plot
plot_ly() %>%
  add_markers(
    x = d$X2, y = d$X1, z = d$Y,
    marker = list(color = "blue", size = 5),
    name = "Data Points"
  ) %>%
  add_surface(
    x = X1_grid, y = X2_grid, z = plane_matrix,
    colorscale = list(c(0, 1), c("red", "pink")),
    showscale = FALSE,
    opacity = 0.7,
    name = "Fitted Plane"
  ) %>%
  plotly::layout(
    scene = list(
      xaxis = list(title = "X1"),
      yaxis = list(title = "X2"),
      zaxis = list(title = "Y")
    )
  )

```

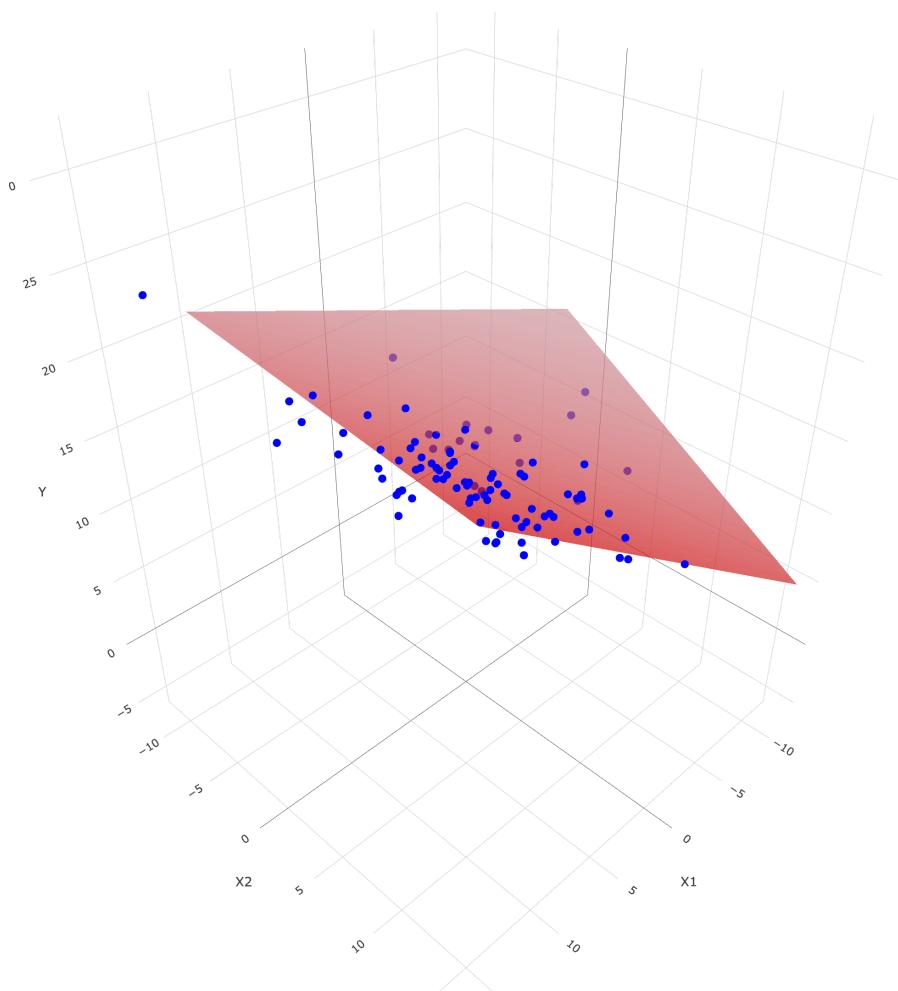
```

),
  title = "Interactive 3D Scatterplot with Fitted Plane"
)

```

```
## file:///private/var/folders/pm/jd6n6gj10371_bml1gh8sc5w0000gn/T/Rtmpj9F7oh/file1574dcddb0d5/w
```

Interactive 3D Scatterplot with Fitted Plane



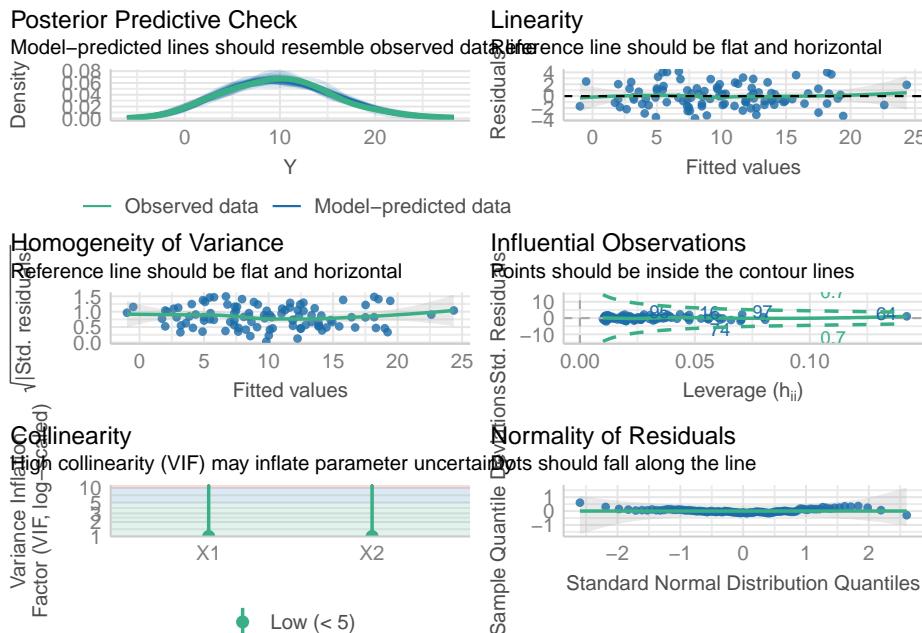
This is, of course, a very idealized situation. There is no curvature in the plane, no interaction, no outliers, no heteroscedasticity. It's the simplest case of multiple regression with 2 predictors. Reality is - usually - more complicated.

Let's look at summary output and check model assumptions:

```
summary(m4.3)
```

```
## 
## Call:
## lm(formula = Y ~ X1 + X2, data = d)
## 
## Residuals:
##     Min      1Q  Median      3Q     Max 
## -3.7460 -1.3215 -0.2489  1.2427  4.1597 
## 
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)    
## (Intercept) 10.27013   0.19228  53.41 <2e-16 ***
## X1          0.44673   0.04195  10.65 <2e-16 ***
## X2          1.00952   0.03960  25.49 <2e-16 ***  
## --- 
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
## 
## Residual standard error: 1.903 on 97 degrees of freedom
## Multiple R-squared:  0.8839, Adjusted R-squared:  0.8815 
## F-statistic: 369.1 on 2 and 97 DF,  p-value: < 2.2e-16
```

```
check_model(m4.3)
```



We could repeat this simulation to get a feeling for the variability. The posterior predictive checks look nice. In this case, we *know* that the model is true and with this knowledge we can assess the diagnostic plots in front of us.

4.2.2.1 Adding variables to the model and why

This is a very complex question. We will go into it in later chapters and the next course (Methodenvertiefung). At this point we can say this: Depending on the goal at hand (prediction or explanation), we add variables to the model and probably use other models apart from linear regression. Prediction seems to be easier than explanation. For instance, within linear models and just a handful of predictors, one can even brute force the problem by searching through all subsets of predictors. If that is not possible, one could use clever algorithms, like best subset selection.

- What is **not** a good idea is to throw all variables into the model and hope for the best.
- What is also not a good idea is to select variables depending on the p -values of the coefficients.
- **Leaving variables out**, that are important, can lead to biased estimates of the coefficients (omitted variable bias).
- Importantly, also **adding variables** can hurt conclusions from the model (see Statistical Rethinking 6.2).

4.2.3 Interaction Term $X_1 \times X_2$

I recommend reading the excellent explanations about interactions in John Kruschke's book Doing Bayesian Data Analysis, 15.2.2 und 15.2.3. Peter Westfall also has a nice explanation in his book in section 9.3.

Our statistical model is now:

$$\begin{aligned} Y_i &= \beta_0 + \beta_1 X_{1i} + \beta_2 X_{2i} + \beta_3 \mathbf{X}_{1i} \times \mathbf{X}_{2i} + \varepsilon_i \\ \varepsilon_i &\sim N(0, \sigma^2) \\ \mathbb{E}(Y_i | X_1 = x_1; X_2 = x_2) &= \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \beta_3 x_1 \times x_2 \\ i &= 1 \dots n \end{aligned}$$

for example:

$$\begin{aligned} Y_i &= 10 + 0.5 \cdot X_{1i} + 1 \cdot X_{2i} + 0.89 \cdot X_{1i} \times X_{2i} + \varepsilon_i \\ \varepsilon_i &\sim N(0, 5) \end{aligned}$$

$$\begin{aligned}\mathbb{E}(Y_i|X_1 = x_1; X_2 = x_2) &= 10 + 0.5x_1 + 1x_2 + 0.89x_1 \times x_2 \\ i &= 1 \dots n\end{aligned}$$

The second equation states that the conditional expectation of Y_i given $X_1 = x_1$ and $X_2 = x_2$ is a function of x_1 and x_2 and their interaction $x_1 \times x_2$. We are in a different situation now. Set for instance x_2 to a certain value, say $x_2 = 7$. Then the relationship (in expectation) between Y and X_1 is:

$$\begin{aligned}\mathbb{E}(Y_i|X_1 = x_1; X_2 = 7) &= 10 + 0.5x_1 + 1 \cdot 7 + 0.89x_1 \cdot 7 \\ \mathbb{E}(Y_i|X_1 = x_1; X_2 = 7) &= 10 + (0.5 + 0.89 \cdot 7) \cdot x_1 + 1 \cdot 7\end{aligned}$$

Depending on the value of x_2 , the *effect* of X_1 on Y changes. Hence, X_2 **modifies** the relationship between X_1 and Y , or stated otherwise, X_1 and X_2 **interact** with respect to Y . Remember, the word *effect* is used in a strictly technical/statistical sense and **not in a causal** sense. It does not mean that if we *do* change X_1 by one unit, Y will also change in an experiment. We are purely describing the relationship in an associative way. We will probably touch causality in a later chapter. Bayesian statistics and causal inference are gaining popularity. Hence, we should try to keep up.

Let's draw 100 points from this model, fit the model and add the plane (see also exercise 4):

```
set.seed(123)
n <- 100
X1 <- rnorm(n, 0, 5)
X2 <- rnorm(n, 0, 5)
Y <- 10 + 0.5 * X1 + 1 * X2 + 0.89 * X1 * X2 + rnorm(n, 0, 5)
d <- data.frame(X1 = X1, X2 = X2, Y = Y)

# Fit the model
m4.4 <- lm(Y ~ X1 * X2, data = d)
summary(m4.4)
```

```
##
## Call:
## lm(formula = Y ~ X1 * X2, data = d)
##
## Residuals:
##    Min     1Q Median     3Q    Max
## -9.360 -3.389 -0.543  2.949 11.583
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)
```

```

## (Intercept) 10.70491   0.47888  22.354 < 2e-16 ***
## X1          0.40719   0.10834  3.759 0.000293 ***
## X2          1.03434   0.09881  10.468 < 2e-16 ***
## X1:X2      0.92182   0.02290  40.257 < 2e-16 ***
## ---
## Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 4.734 on 96 degrees of freedom
## Multiple R-squared:  0.9476, Adjusted R-squared:  0.9459
## F-statistic: 578.1 on 3 and 96 DF,  p-value: < 2.2e-16

# Create a grid for the plane
X1_grid <- seq(min(d$X1), max(d$X1), length.out = 20)
X2_grid <- seq(min(d$X2), max(d$X2), length.out = 20)
grid <- expand.grid(X1 = X1_grid, X2 = X2_grid)

# Predict the values for the grid
grid$Y <- predict(m4.4, newdata = grid)

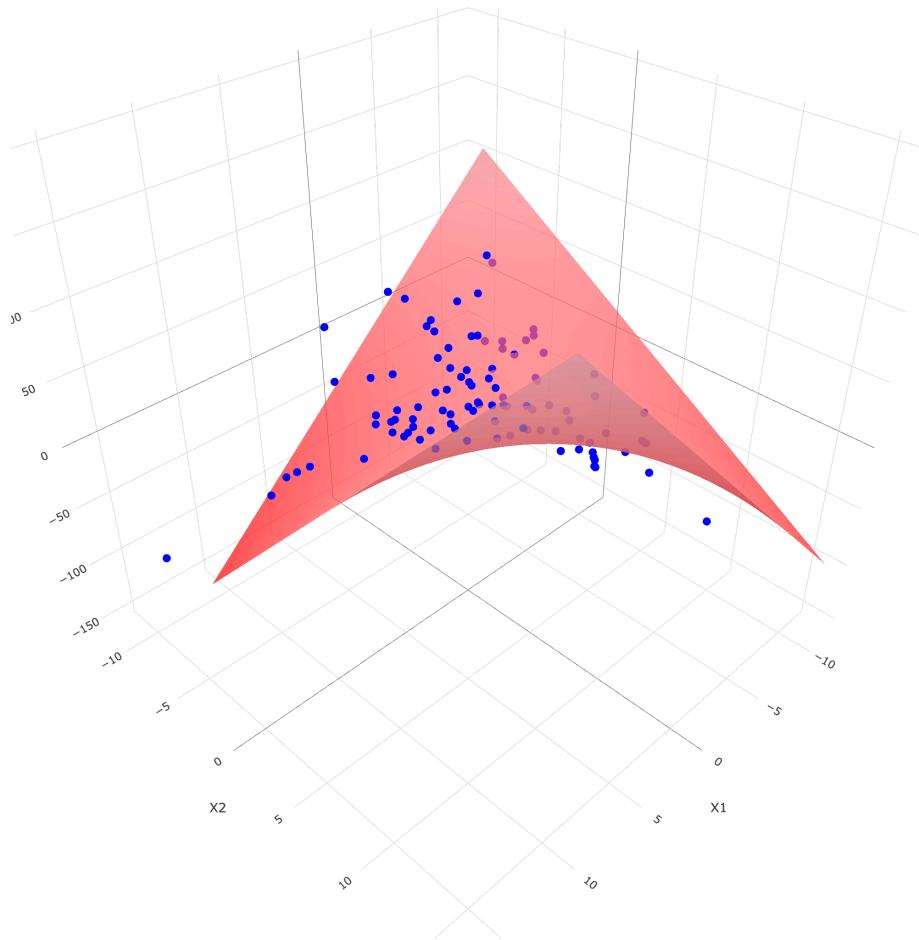
# Convert the grid into a matrix for the plane
plane_matrix <- matrix(grid$Y, nrow = length(X1_grid), ncol = length(X2_grid))

# Create the interactive 3D plot
plot_ly() %>%
  add_markers(
    x = d$X2, y = d$X1, z = d$Y,
    marker = list(color = "blue", size = 5),
    name = "Data Points"
  ) %>%
  add_surface(
    x = X1_grid, y = X2_grid, z = plane_matrix,
    colorscale = list(c(0, 1), c("red", "pink")),
    showscale = FALSE,
    opacity = 0.7,
    name = "Fitted Plane"
  ) %>%
  plotly::layout(
    scene = list(
      xaxis = list(title = "X1"),
      yaxis = list(title = "X2"),
      zaxis = list(title = "Y")
    ),
    title = "Interactive 3D Scatterplot with Fitted Plane"
  )

## file:///private/var/folders/pm/jd6n6gj10371_bml1gh8sc5w0000gn/T/Rtmpj9F7oh/file1574d268c72c/w

```

Interactive 3D Scatterplot with Fitted Plane



The term $X1 * X2$ is a shortcut for $X1 + X2 + X1:X2$ where $X1:X2$ is the interaction term. R automatically includes the main effects of the predictors when an interaction term is included. The true but usually unknown β s are estimated quite precisely.

4.2.3.1 Formal test for interaction

We could apply a formal test for the interaction term by model comparison. The command `anova(., .)` would compare the two models and test if the change in the residual sum of squares is statistically interesting.

```
m4.5 <- lm(Y ~ X1 + X2, data = d) # without interaction
anova(m4.5, m4.4)

## Analysis of Variance Table
##
## Model 1: Y ~ X1 + X2
## Model 2: Y ~ X1 * X2
##   Res.Df   RSS Df Sum of Sq    F    Pr(>F)
## 1     97 38468
## 2     96 2151  1    36316 1620.6 < 2.2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

One can show that the following test statistic is F distributed under the null hypothesis:

$$F = \frac{(RSS_{\text{Model 1}} - RSS_{\text{Model 2}}) / (df_{\text{Model 1}} - df_{\text{Model 2}})}{RSS_{\text{Model 2}} / df_{\text{Model 2}}}$$

where RSS is the residual sum of squares, df are the degrees of freedom of the residual sum of squares for both models.

The output of the `anova` command shows us the residual degrees of freedom (`Res.Df`) of both models, the residual sum of squares errors of both models (`RSS`), the sum of squared errors between model 1 and model 2 (`Sum of Sq`), the value of the F-statistic and the p -value for the hypothesis, that the coefficient for the interaction term is zero ($\beta_3 = 0$). Model 1 RSS has 97 degrees of freedom, since we have 100 data points and 3 parameters to estimate ($\beta_0, \beta_1, \beta_2$). Model 2 has 96 degrees of freedom, since we have 100 data points and 4 parameters to estimate ($\beta_0, \beta_1, \beta_2, \beta_3$).

Let's verify the value of the F statistic:

```
RSS_model1 <- sum(residuals(m4.5)^2)
RSS_model2 <- sum(residuals(m4.4)^2)
df_model1 <- n - length(coef(m4.5))
df_model2 <- n - length(coef(m4.4))
F <- ((RSS_model1 - RSS_model2) / (df_model1 - df_model2)) / (RSS_model2 / df_model2)
F

## [1] 1620.606

# Sum of Sq
RSS_model1 - RSS_model2
```

```
## [1] 36316.28
```

In the numerator of the F statistic, we have the change in the residual sum of squares (from the small (model 1) model to the larger one (model 2), **Sum of Sq**) per additional parameter in the model (one additional parameter β_3).

In the denominator, we have the residual sum of squares per residual degree of freedom of the larger model (model 2). Hence, in the numerator we have the information on how much better we get with respect to the number of variables added, and in the denominator we have information on how good the full model is with respect to its degrees of freedom.

The p -value is the probability of observing a value of the F statistic as extreme or more extreme than the one we observed, given that the null hypothesis is true. Here, the p -value is extremely small. So, statistically we would see an improvement in RSS which is not explainable by chance alone. But **let's be careful with p -values** and especially with fixed cutoff values for α , which we will **never** use in this script. Even for a rather small effect β_3 , we would reject the null hypothesis, if only the sample size is large enough. Since a very small effect relative to β_1 and β_2 would probably not be of practical interest, one should be careful with looking at p -values alone. For instance, in Richard McElreath's book Statistical Rethinking, there are no p -values at all. I like that.

If you again look at the comparison of the RSS between the two models, you would immediately see that the model with the interaction term is better (at least with respect to this metric). The difference is huge. We have already mentioned in the context of R^2 not to overinterpret such metric, because RSS is monotonically decreasing with number of variables added and reaches zero when the number of variables equals the number of data points (see exercise 3).

4.2.4 Using an interaction plot to see a potential interaction

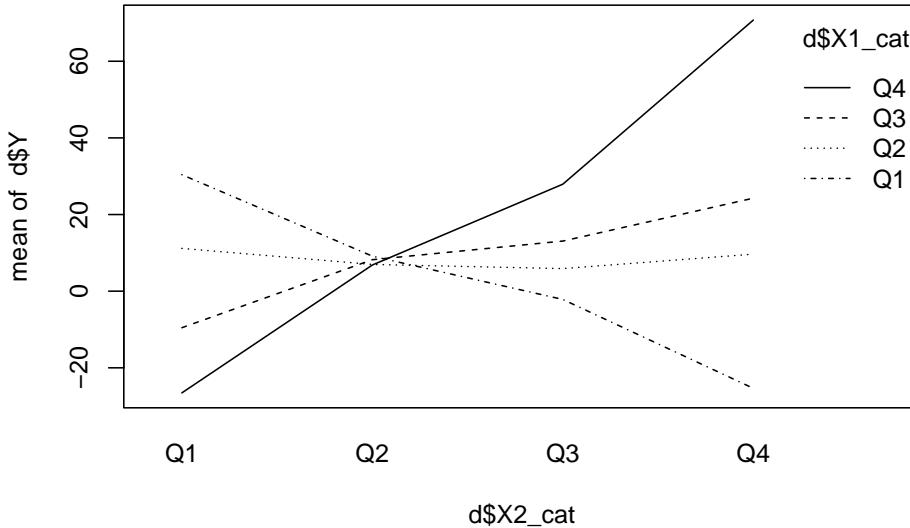
Chronologically before we include an interaction term in the model, we can use an interaction plot to see if there is a potential interaction between the predictors. We can just create a categorical predictor out of the continuous predictors. We just categorize the predictors into quartiles and plot the means of the dependent variable (Y). If the lines are parallel, there is no interaction. If the lines are not parallel, there might be an interaction.

```
n <- 100
X1 <- rnorm(n, 0, 5)
X2 <- rnorm(n, 0, 5)
Y <- 10 + 0.5 * X1 + 1 * X2 + 0.89 * X1 * X2 + rnorm(n, 0, 5)
d <- data.frame(X1 = X1, X2 = X2, Y = Y)
```

```
# Create categorical variables based on quartiles
d$X2_cat <- cut(d$X2,
                  breaks = quantile(d$X2, probs = c(0, 0.25, 0.5, 0.75, 1), na.rm = TRUE),
                  include.lowest = TRUE,
                  labels = c("Q1", "Q2", "Q3", "Q4"))

d$X1_cat <- cut(d$X1,
                  breaks = quantile(d$X1, probs = c(0, 0.25, 0.5, 0.75, 1), na.rm = TRUE),
                  include.lowest = TRUE,
                  labels = c("Q1", "Q2", "Q3", "Q4"))

# Create the interaction plot
interaction.plot(d$X2_cat, d$X1_cat, d$Y)
```



There seems to be an interaction of the predictors with respect to Y . The lines are not parallel. If there was no interaction, the change in Y with respect to X_1 would be the same for all levels of X_2 . This seems not to be the case here. See exercise 5.

If we had one or both predictors already categorical, we would not have to discretize them before.

4.2.5 Simpsons Paradox and what happens when you just throw variables into multiple regression

The Simpsons paradox is a phenomenon, in which a trend appears in several different groups of data but disappears or reverses when these groups are combined. I agree with the criticism that this is not really a paradox but a failure

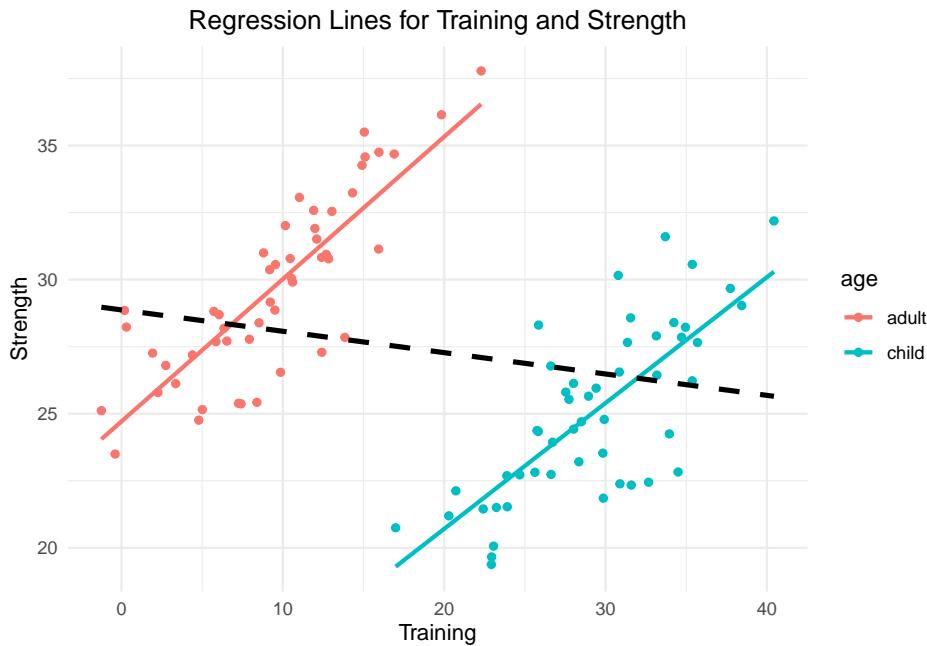
to consider confounding variables adequately. Let's quickly invent an example. We are interested in the relationship hours of muscle training and strength (not based on evidence) in children vs. adults. Within both groups there will be an increasing relationship. The more training, the more muscle strength. But if we combine the groups, we will see a decreasing relationship.

```
library(tidyverse)
n <- 100
age <- c(rep("child", n/2), rep("adult", n/2))
training <- c(rnorm(n/2, 0, 5) + 30, rnorm(n/2, 0, 5) + 10)
strength <- c(
  10 + 0.5 * training[1:(n/2)] + rnorm(n/2, 0, 2), # For children
  25 + 0.5 * training[(n/2 + 1):n] + rnorm(n/2, 0, 2) # For adults
)

d <- data.frame(age = age, training = training, strength = strength)

ggplot(d, aes(x = training, y = strength, color = age)) +
  geom_point() +
  geom_smooth(method = "lm", se = FALSE) + # Group-specific regression lines
  geom_smooth(data = d, aes(x = training, y = strength),
              method = "lm", se = FALSE, color = "black", linetype = "dashed", linewidth = 1)
  labs(title = "Regression Lines for Training and Strength",
       x = "Training",
       y = "Strength") +
  theme_minimal() +
  theme(plot.title = element_text(hjust = 0.5))

## `geom_smooth()` using formula = 'y ~ x'
## `geom_smooth()` using formula = 'y ~ x'
```



- **Group-Specific Trends:**

- In the group of **children** (blue line), strength **increases** with training, as indicated by the positive slope of the regression line.
- Similarly, in the group of **adults** (red line), strength also **increases** with training.

- **Overall Trend:**

- When both groups are combined, the overall regression line (black, dashed) shows a **negative slope**, suggesting that **strength decreases** with training.
- This overall trend is opposite to the trends observed within the individual groups.

- **Why Does This Happen?**

- This paradox occurs because the relationship between the grouping variable (**age**) and the independent variable (**training**) creates a confounding effect.
- In this case:
 - * **Children** tend to have higher training values overall, while **adults** tend to have lower training values.
 - * The overall regression line reflects the imbalance between these groups rather than the true within-group relationships.

This (fictitious) example illustrates the importance of considering group-level differences when analyzing data. In multiple regression this might not be so straightforward, since we cannot just plot the data and see the effect as in our toy example. We will discuss this later.

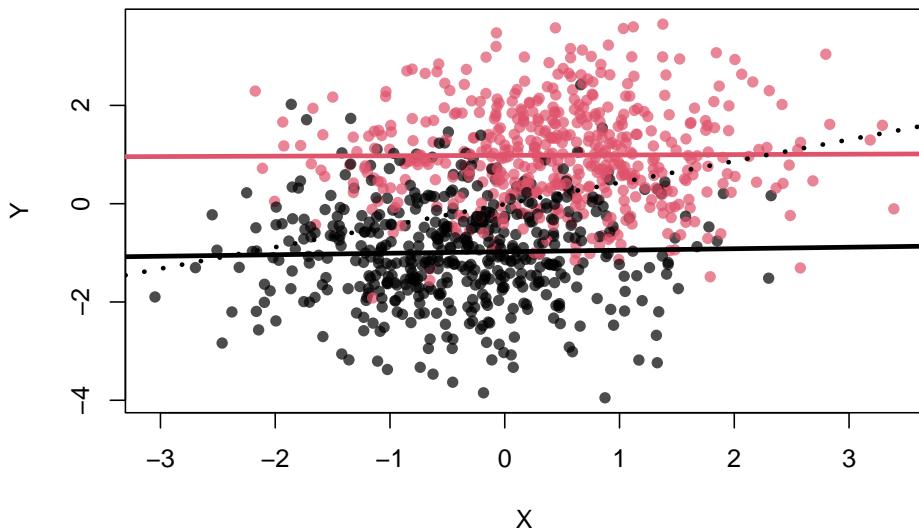
Richard McElreath also has a nice example for this on github and explains this in a video.

```
library(rethinking)

# d-separation plots
a <- 0.7
cols <- c( col.alpha(1,a) , col.alpha(2,a) )

# pipe
# X -> Z -> Y
N <- 1000
X <- rnorm(N)
Z <- rbern(N,inv_logit(X))
Y <- rnorm(N,(2*Z-1))

plot( X , Y , col=cols[Z+1] , pch=16 )
abline(lm(Y[Z==1] ~ X[Z==1]),col=2,lwd=3)
abline(lm(Y[Z==0] ~ X[Z==0]),col=1,lwd=3)
abline(lm(Y~X),lwd=3,lty=3)
```



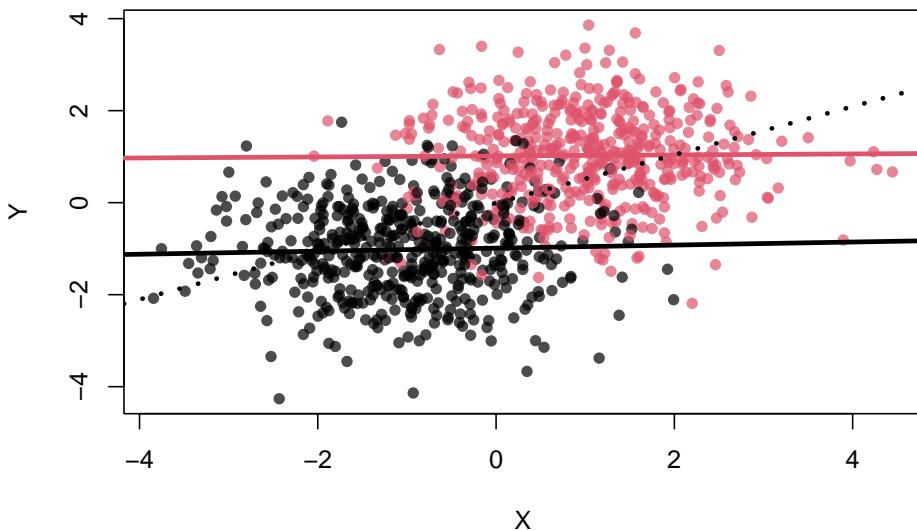
```
# fork
# X <- Z -> Y
```

```

N <- 1000
Z <- rbern(N)
X <- rnorm(N,2*Z-1)
Y <- rnorm(N,(2*Z-1))

plot( X , Y , col=cols[Z+1] , pch=16 )
abline(lm(Y[Z==1]~X[Z==1]),col=2,lwd=3)
abline(lm(Y[Z==0]~X[Z==0]),col=1,lwd=3)
abline(lm(Y~X),lwd=3,lty=3)

```

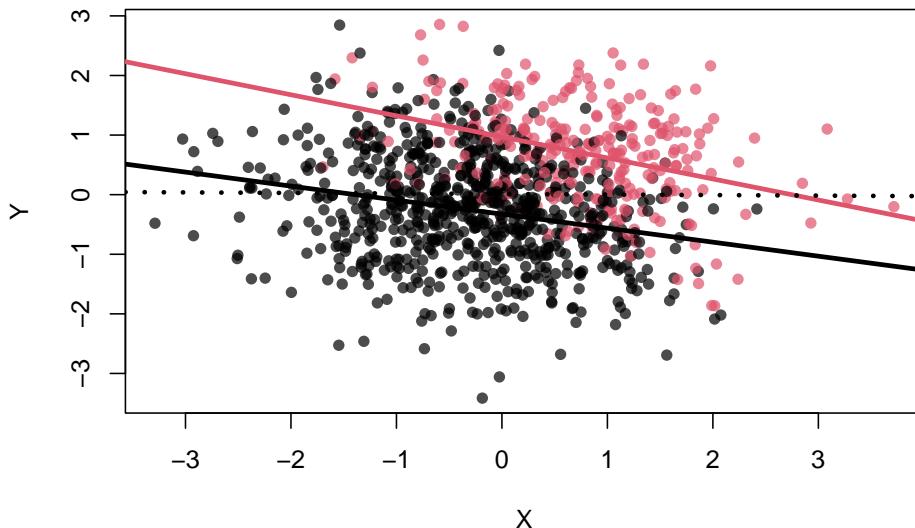


```

# collider
# X -> Z <- Y
N <- 1000
X <- rnorm(N)
Y <- rnorm(N)
Z <- rbern(N,inv_logit(2*X+2*Y-2))

plot( X , Y , col=cols[Z+1] , pch=16 )
abline(lm(Y[Z==1]~X[Z==1]),col=2,lwd=3)
abline(lm(Y[Z==0]~X[Z==0]),col=1,lwd=3)
abline(lm(Y~X),lwd=3,lty=3)

```



- Pipe: In this setting X is associated with Z and Z is associated with Y . X and Y are also associated. But if we condition on Z , the association between X and Y disappears. This is the pipe. This means, if we **know** the value of Z , X does not give us any information about Y . This can be seen in the first plot: Once we are within $Z = 0$ or $Z = 1$, X does not give us any information about Y , i.e., the point cloud is horizontal and there is no correlation.

Or in the framework of Simpons paradox:

```
# pipe
N <- 1000
X <- rnorm(N)
Z <- rbern(N, inv_logit(X))
Y <- rnorm(N, (2*Z-1))
mod1 <- lm(Y ~ X)
summary(mod1)

##
## Call:
## lm(formula = Y ~ X)
##
## Residuals:
##      Min      1Q  Median      3Q     Max 
## -4.1686 -0.9507 -0.0105  0.9292  3.7899 
## 
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)    
## (Intercept)  0.0000    0.9950   0.000   1.000    
## X           2.0000    0.9950   2.005   0.045 *  
## ---
```

```

## (Intercept) -0.004706   0.042844   -0.11    0.913
## X           0.432429   0.041339   10.46 <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1.355 on 998 degrees of freedom
## Multiple R-squared:  0.09881, Adjusted R-squared:  0.0979
## F-statistic: 109.4 on 1 and 998 DF, p-value: < 2.2e-16

mod2 <- lm(Y ~ X + Z)
summary(mod2)

## 
## Call:
## lm(formula = Y ~ X + Z)
## 
## Residuals:
##      Min       1Q   Median       3Q      Max
## -3.4382 -0.6174 -0.0548  0.6945  3.0373
## 
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)
## (Intercept) -1.017412  0.046595 -21.835 <2e-16 ***
## X            -0.004415  0.033723  -0.131   0.896
## Z            2.055013  0.069907  29.396 <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.9921 on 997 degrees of freedom
## Multiple R-squared:  0.5172, Adjusted R-squared:  0.5163
## F-statistic: 534.1 on 2 and 997 DF, p-value: < 2.2e-16

```

Adding Z to the model makes the coefficient for X disappear.

See exercise 7.

4.3 Exercises

[E] Easy, [M] Medium, [H] Hard

(Some) solutions to exercises can be found in the git-repo here.

4.3.1 [M] Exercise 1

- Fit a model with a cubic term for weight and height of the !Kung San people.
- Add the prediction bands as seen in the book.
- Come up with an explanation for the functional form of this relationship.
- Could there be reasons to for taking a less complicated model (1, 2)?

4.3.2 [E] Exercise 2

Consider the model equations from above where we used polynomial regression to model the relationship between weight and height:

- Draw the model hierarchy for the model.

4.3.3 [H] Exercise 3

Invent a data set (or use the first 3 lines of a previous data set) with 3 observations of Y and X_1, X_2 and X_3 . You have a data frame with 3 rows and 4 columns.

- Fit a model with Y as the dependent variable and X_1, X_2, X_3 as predictors.
- How big is R^2 ?
- Could you have calculated this without lm and R?

4.3.4 [E] Exercise 4

Go back to the section about the interaction term in the linear model.

- Use the code provided.
- Standardise the predictors. How are the β s changing and what is their interpretation now?
- Change the relative sizes of the true but usually unknown β s. What happens to the estimates and the graph?
- What happens if you change the error term and increase or decrease its variance?

4.3.5 [E] Exercise 5

Draw the interaction plot from the section about the interaction plot for the case when there is no interaction, i.e. $\beta_3 = 0$.

4.3.6 [M] Exercise 6

Go back to the model assumptions checks above.

- Create the same two plots for the simple mean model without predictors, just with the intercept.
- Which model fits the data better according to these posterior predictive checks?

4.3.7 [E] Exercise 7

Go back to the Simpson's paradox section.

- Verify that the coefficient for X disappears when you add Z to the model for the pipe.
- What happens in the case of the collider?

4.3.8 [M] Exercise 8

- Take a data set of your choosing with many columns, say 10 or so, either from the internet or from R (if available)
- Fit a model for an arbitrary outcome, add more and more variables to predict the outcome and verify that the R^2 increases.

4.3.9 [M] Exercise 9

Exponential curve fitting. Go back to the section about adding a transformed predictor in the Frequentist setting. Hint: You can use the `optim` function in R.

- Assume the relationship between weight and height looks like this: $height_i = \alpha + \beta_1 e^{\beta_2 weight_i}$.
- Use R and the least squares method to estimate the parameters α, β_1, β_2 .
- Note that the sum of squared errors is this: $\sum_{i=1}^n (height_i - \alpha - \beta_1 e^{\beta_2 weight_i})^2$.
- What happens if you do not constrain the parameters β_1 and β_2 to be negative?
- Calculate the R^2 for this model.

4.4 TODOS

- `check_model` for 2 and more predictors...
- Violations of model assumptions and how to fix them.
- Clarify what of pipe, collider and fork is relevant for the simposons paradox.
- Add more Bayesian model checks, chapter 6 Gelman.
- Show by simulation what Gelman talks about with significant p values. So I scan the data for significant p values and then simulate data with the same effect size and see how often I get significant p values. Especially the next effect would be probably smaller, especially, if one did p-hacking! Calculate a priori probability for replication (def?).
- Variability of confidence interval borders (draw from X...)
- Logistic Regression, Poisson, ...
- Chapter: Sample size calculations for logistic and multivariate regression, Proportions, ICCs, t.test
- Chapter about Reliability, Validity and ICCs (incl. simulation of what an ICC of 0.9 or so means), but maybe reduced
- causal terror of richards book, fork, pipe, collider
- Angenommen man hat ein masking eines Effekts und der Model fit ist aber gut (keine Voraussetzung verletzt), ist diese Situation möglich?
- What about 2 interactions?
- Which variables should I include in a model and why?
- What about a real life data set?
- What about papers?
- Simulate a cohort of Master thesis with small sample sizes and assume that there is a true but unknown effect.
- AIC, BIC, cross-validation, Model selection (best subset, leaps...), Variable selection